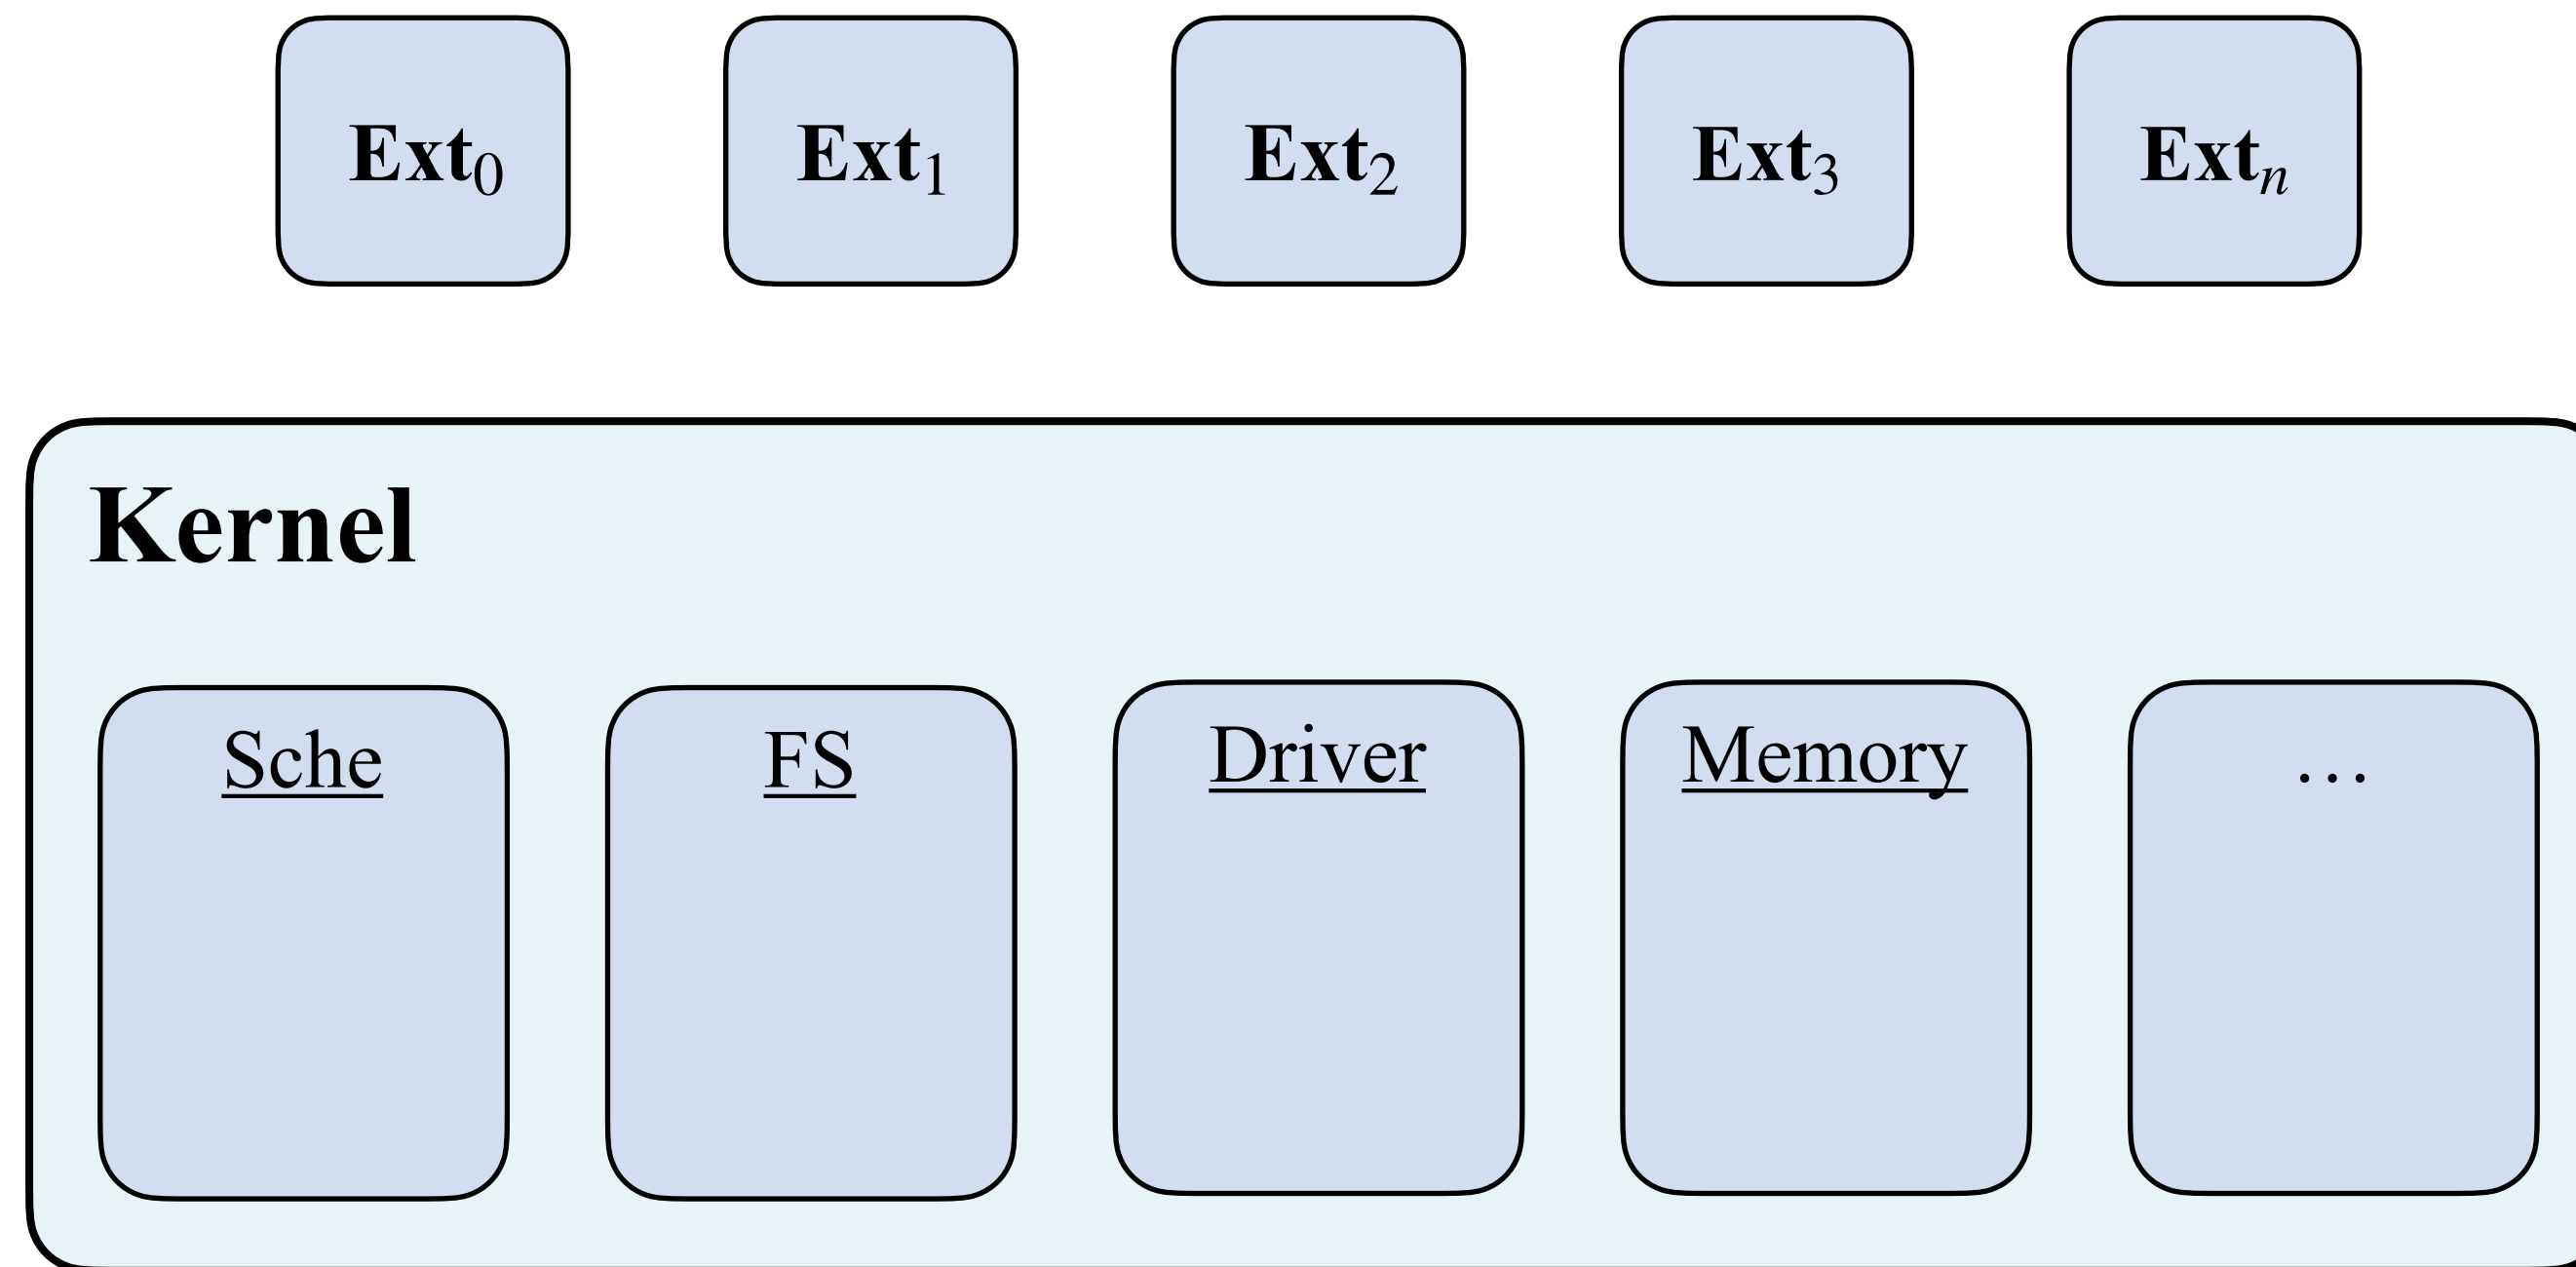




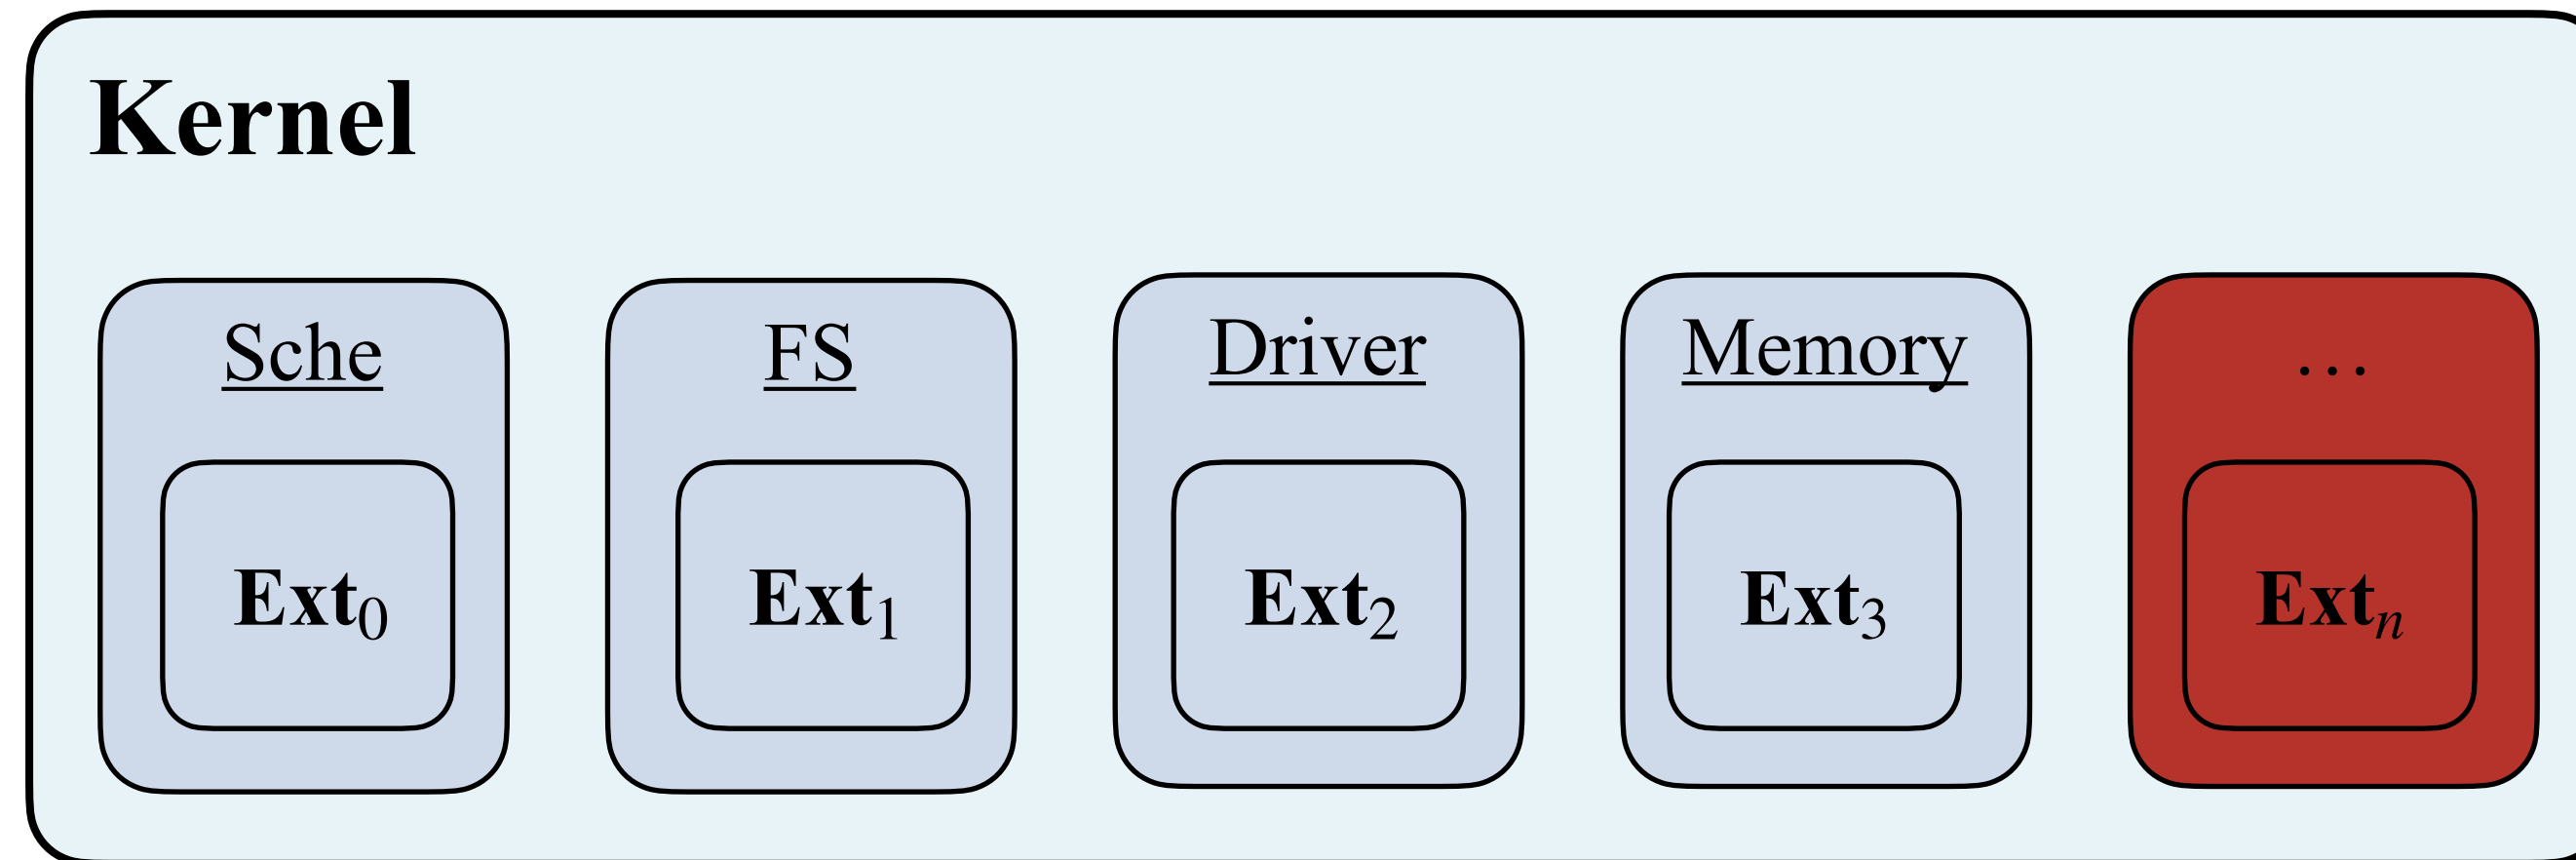
# Prove It to the Kernel: Precise Extension Analysis via Proof-Guided Abstraction Refinement

**Hao Sun, Zhendong Su**  
**ETH Zurich**

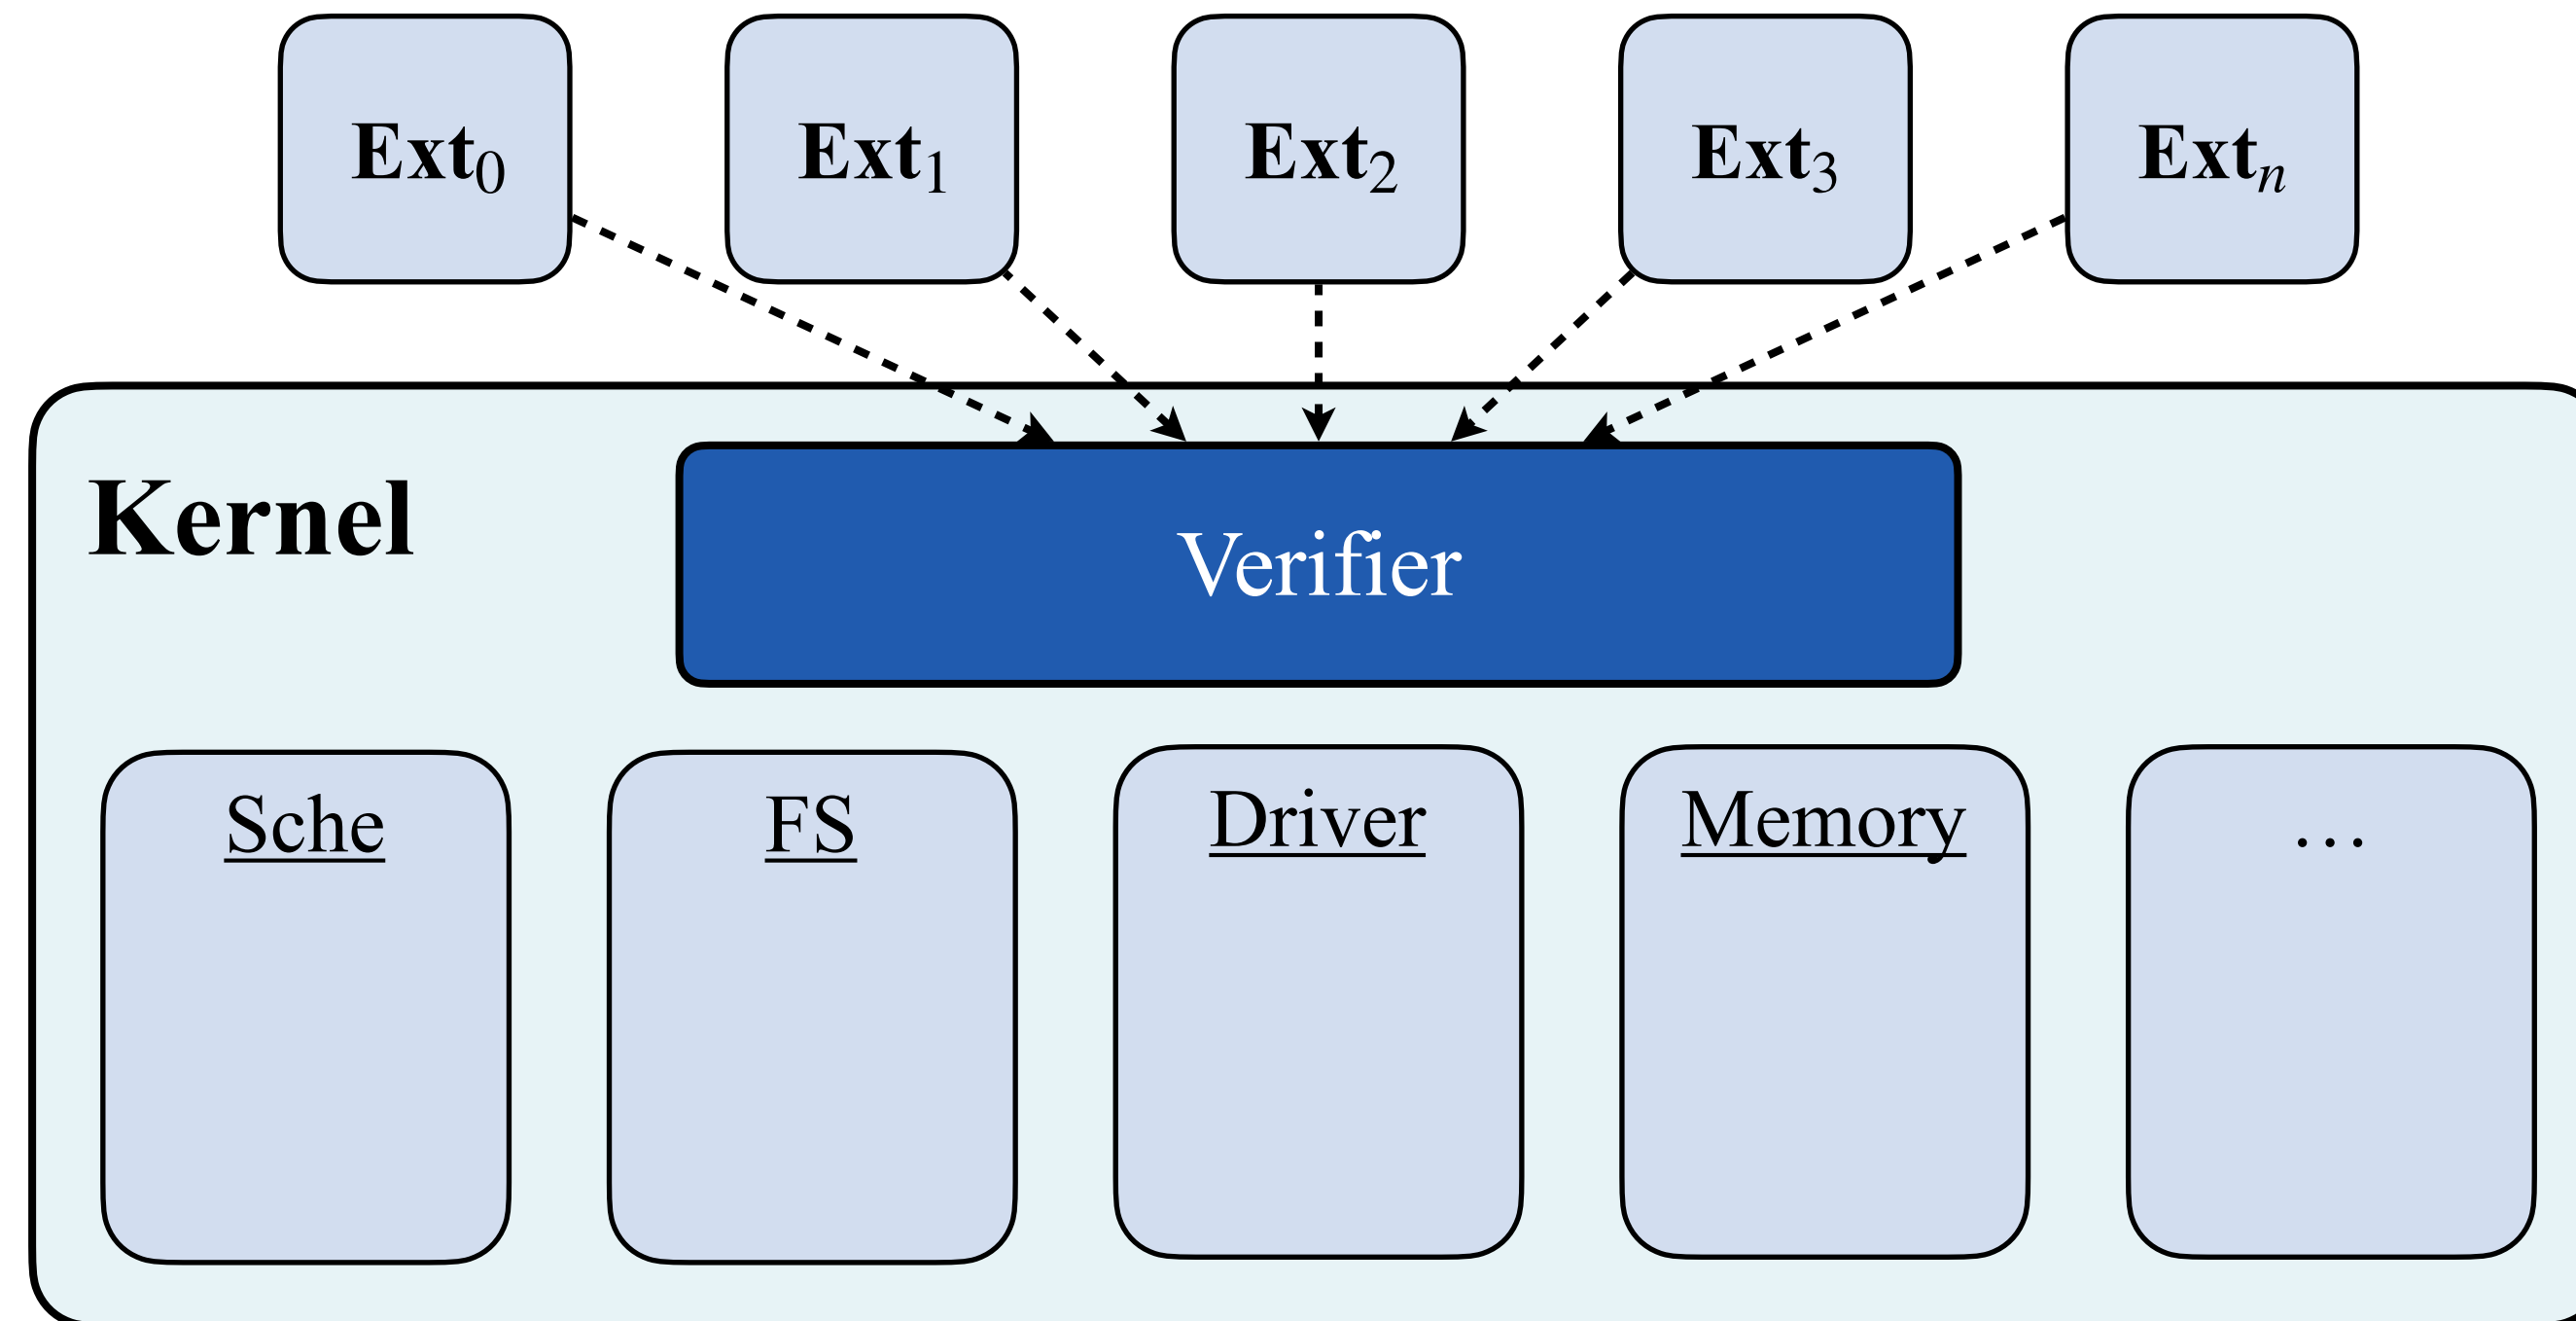
# Extensible Kernel



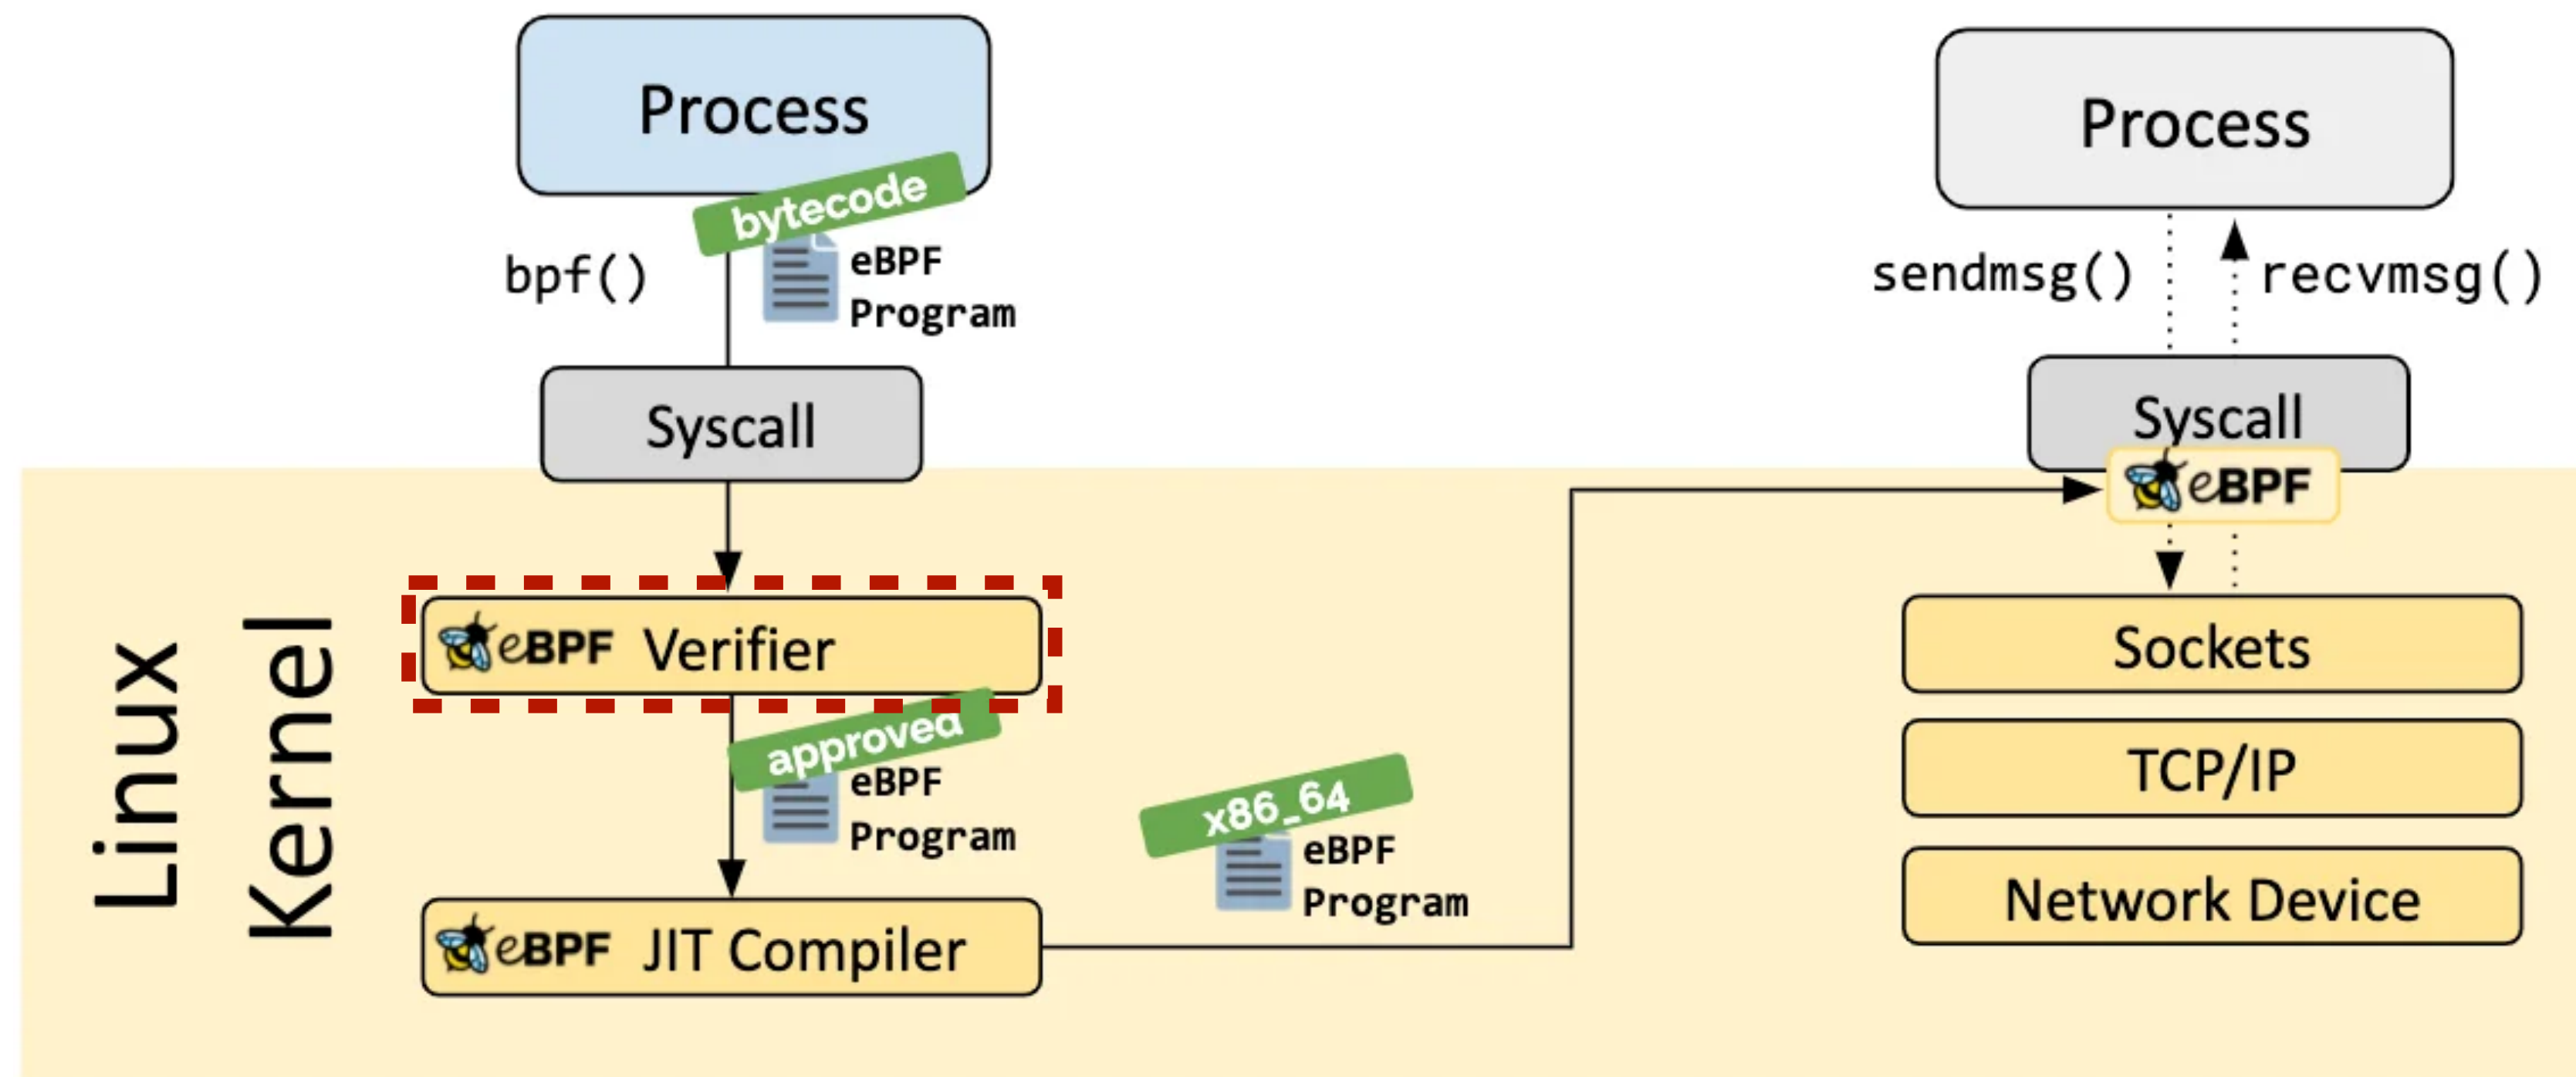
# Extension Safety



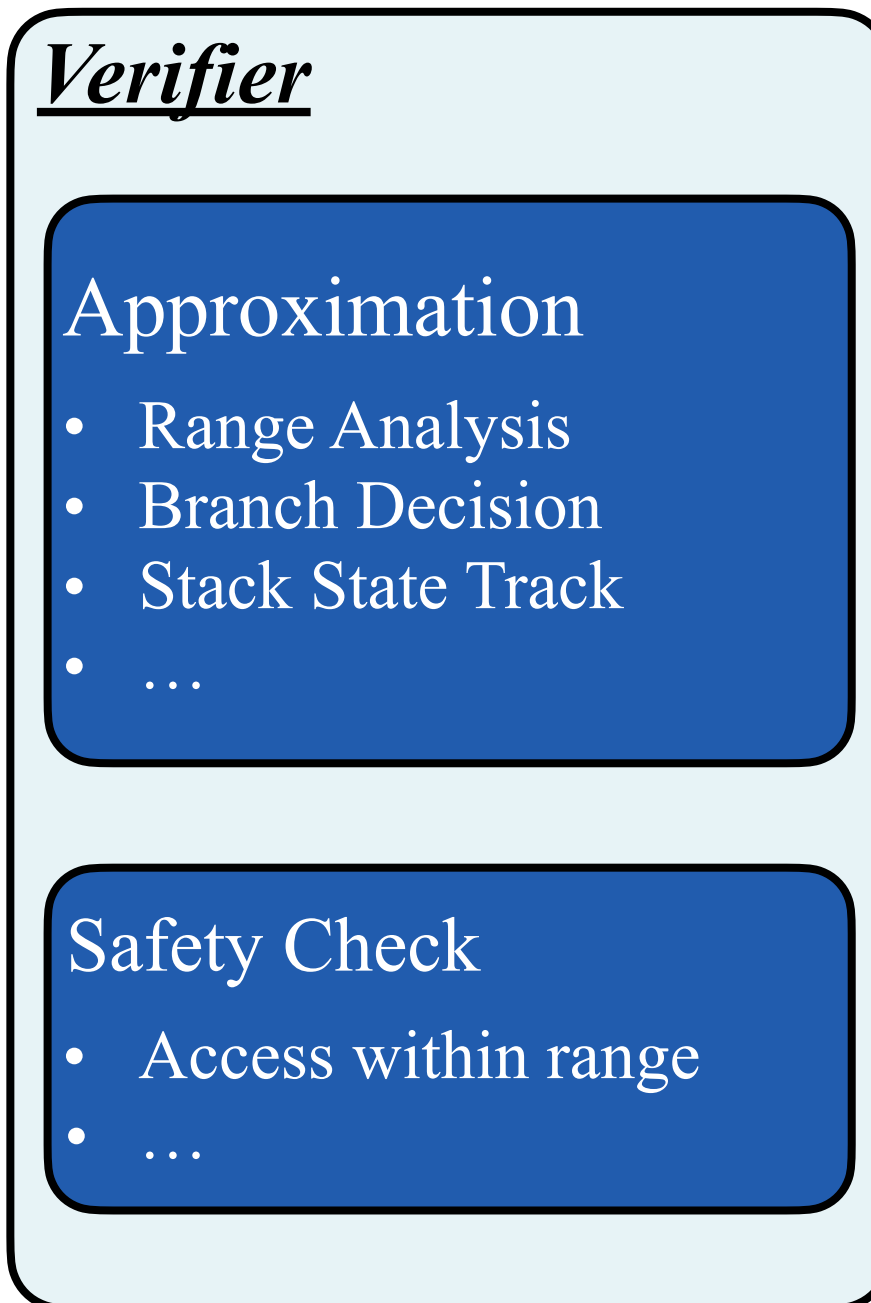
# In-Kernel Extension Analysis



# eBPF



# eBPF Verifier

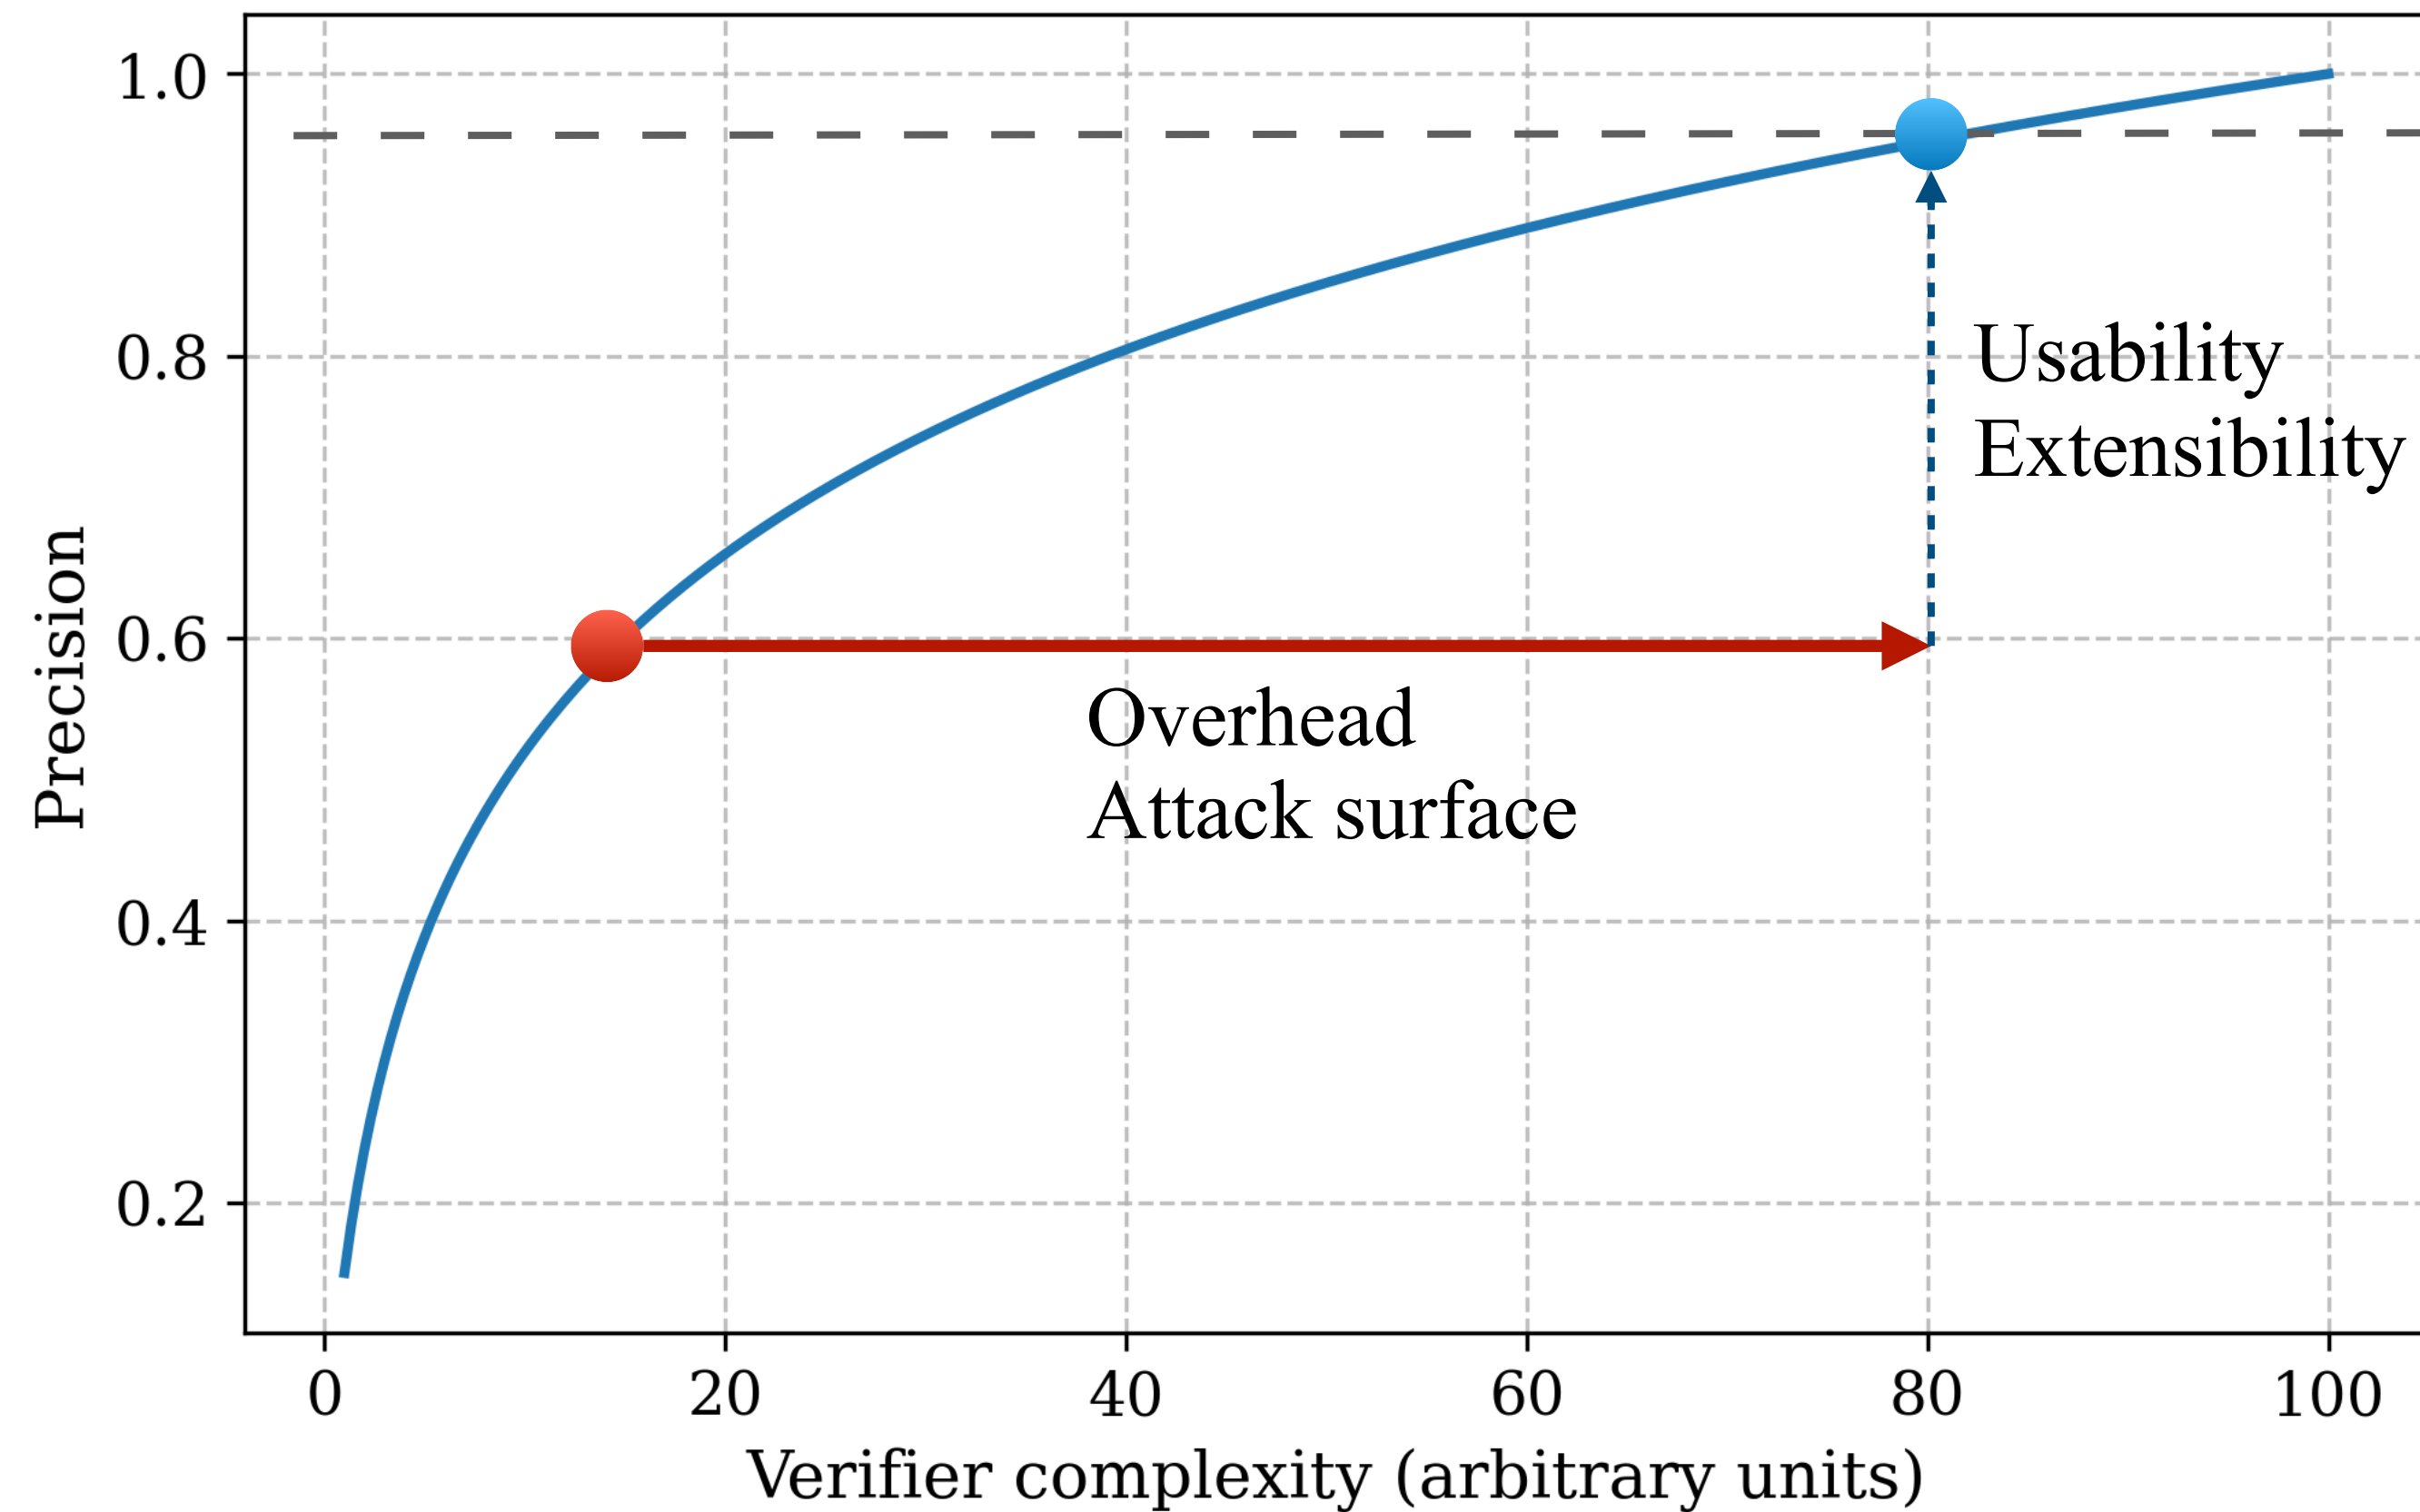



```
0: r1 = map_lookup(...) // r1 = ptr(size=16)
1: r2 &= 0xf             // r2 ∈ [0,15]
2: r2 <=<= 1             // r2 ∈ [0,30]
3: r1 += r2              // r1.off ∈ [0,30]
4: r0 = *(u8*)r1         // unsafe: off+access_sz > mem_sz
```

**Listing 1.** Correct rejection of unsafe extension. The comments illustrate the verifier's analysis.

- Approximate the possible value set of each variable
- Check operation safety against the approximation
- Load the extension only if it passes the verifier

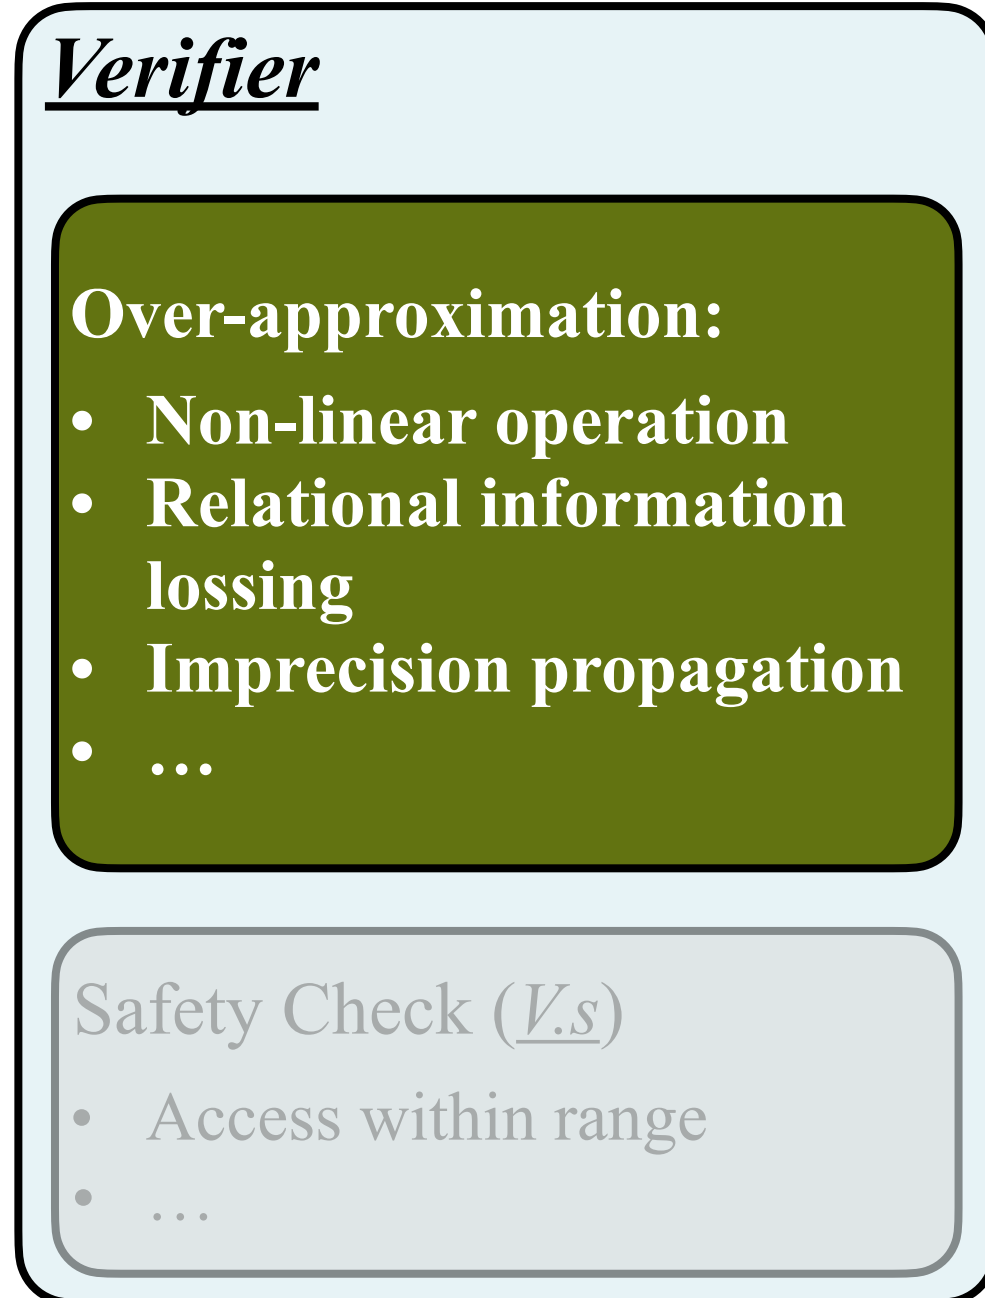
# Precision vs. Complexity



**Key challenge:** achieving **higher precision** while maintaining **low complexity**



# Imprecisions



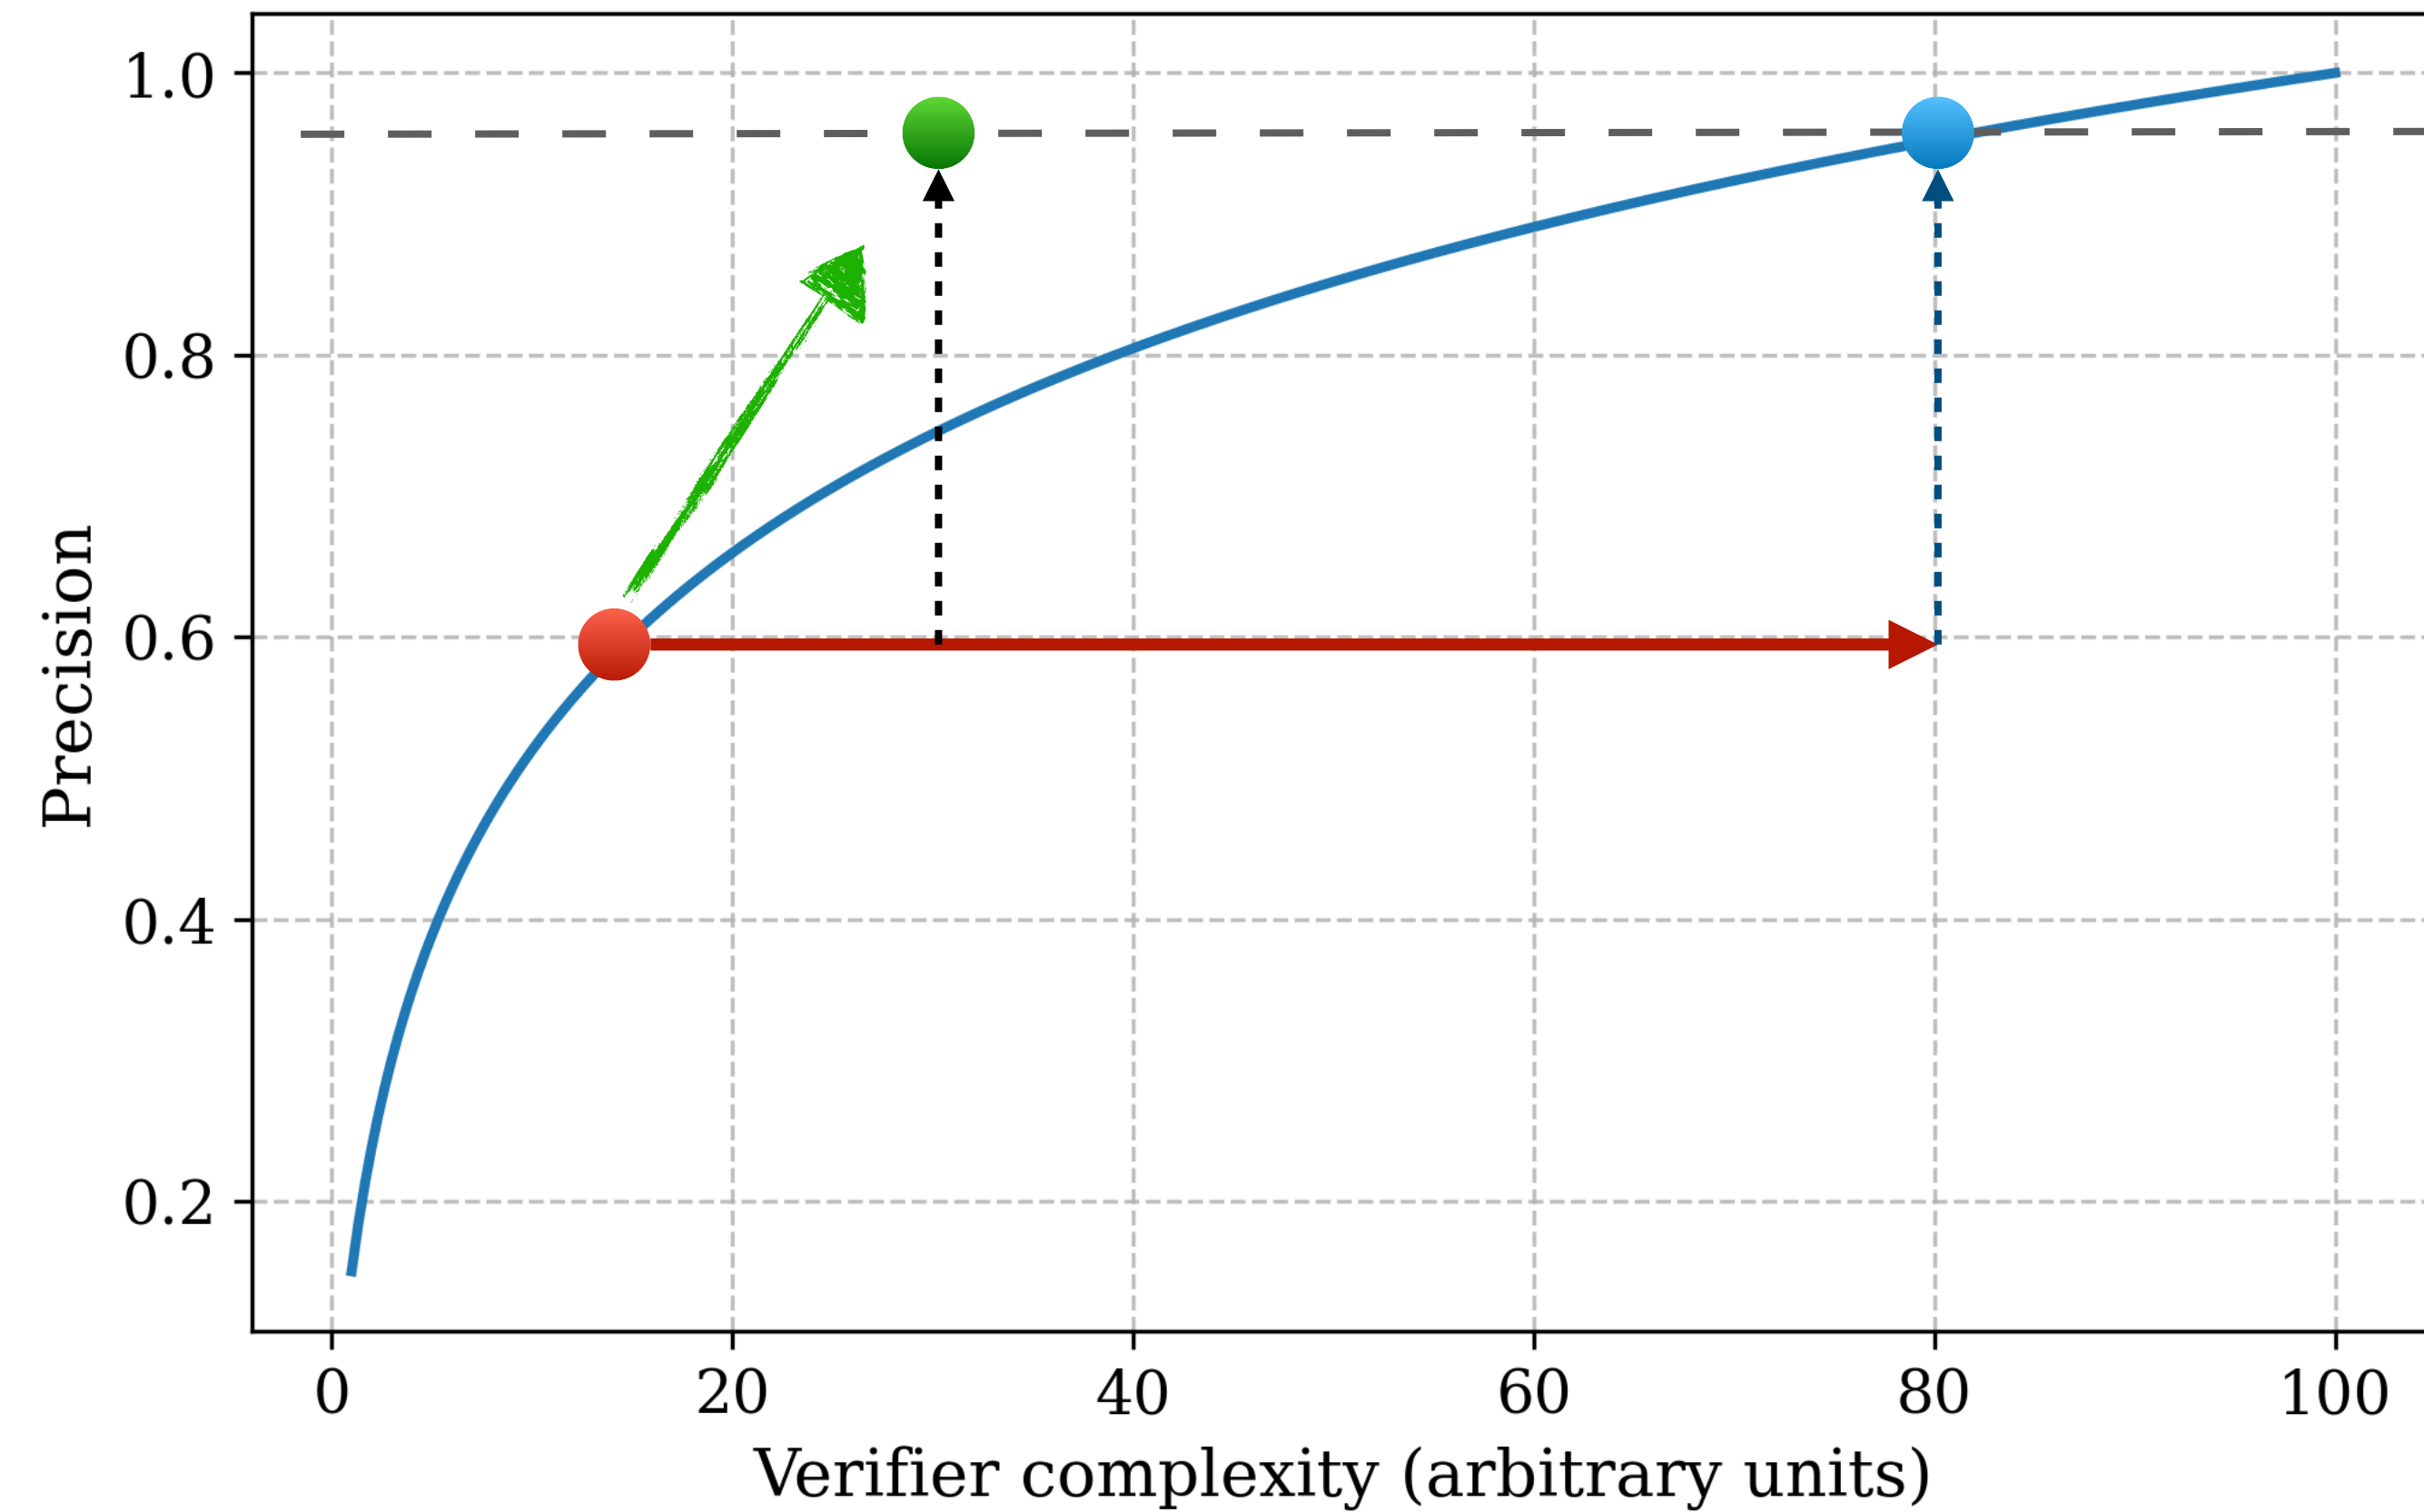

Instruction	Abstraction
1: r1 = map_lookup(...)	$r1 = ptr(\underline{size=16})$
2: r2 &= 0xf	$r2 \in [0, 15]$
3: r1 += r2	$r1.off \in [0, 15]$
4: r3 = 0xf - r2	$r3 \in [0, 15]$
5: r1 += r3	$r1.off \in [0, 30]$
6: r0 = *(u8*)r1	Unsafe, rejected..

```
// 256 KiB per cpu core, of which 128 KiB is usable as
// we have to bound each new variable-length field to
// start at no more than half the size of the buffer to
// make the verifier happy.
#define EVENT_BUFFER_SIZE (1 << 18)
#define EVENT_BUFFER_SIZE_HALF (EVENT_BUFFER_SIZE >> 1)
```

**Listing 3.** Workaround to mitigate the verifier's imprecision.

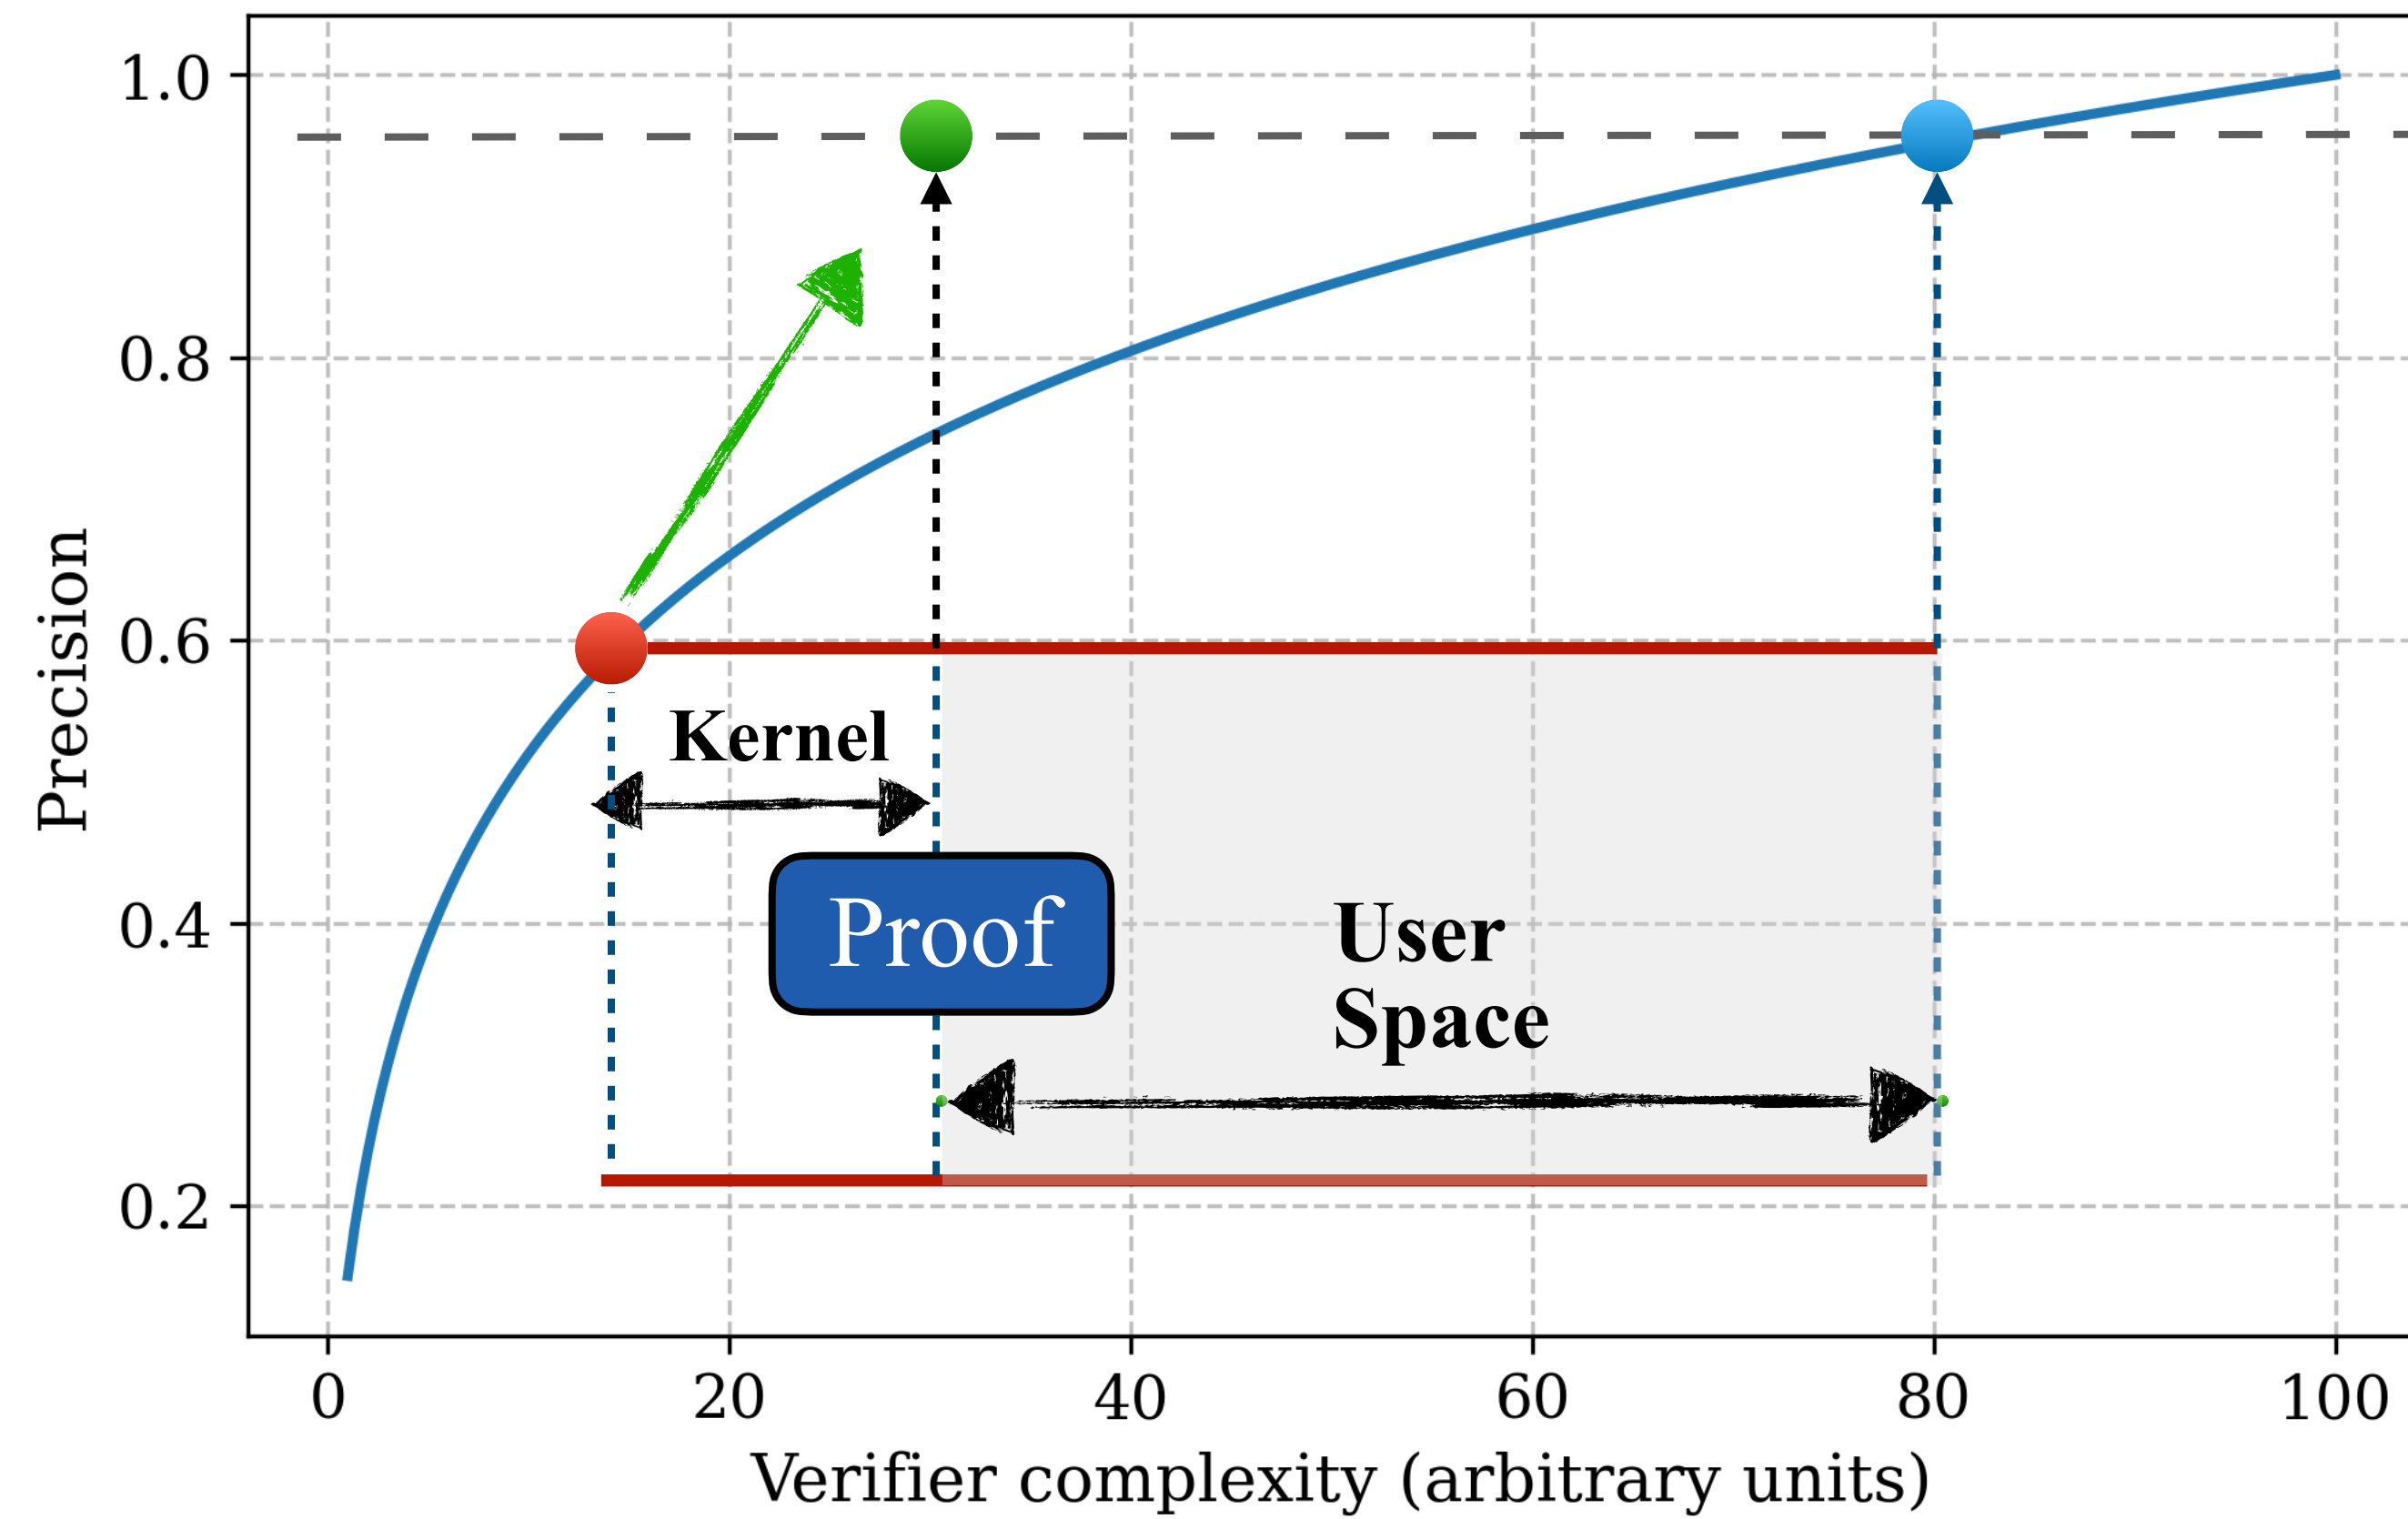


# Our Goal



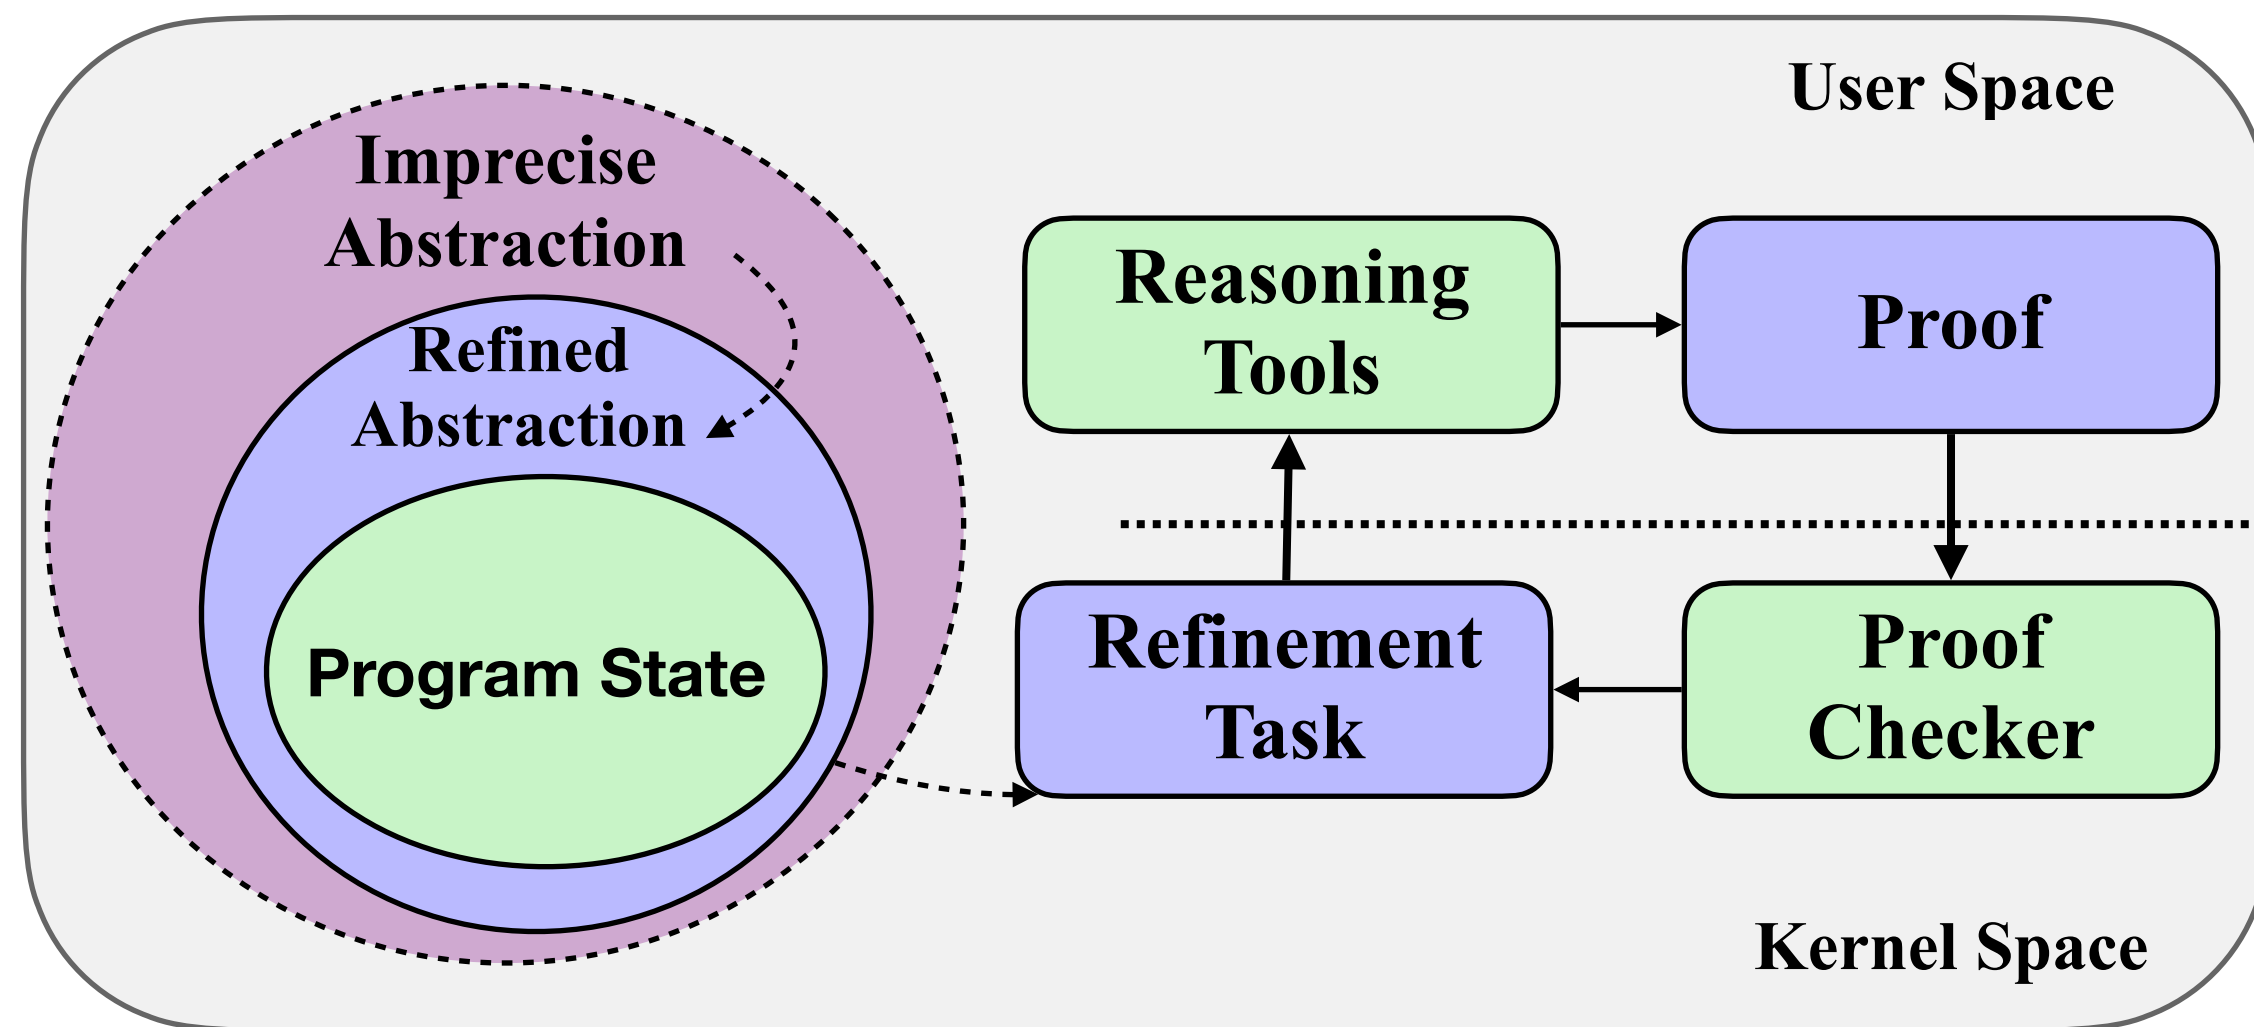
**Goal:** significantly enhance the precision with **linear-time** kernel space complexity.

# Key Idea




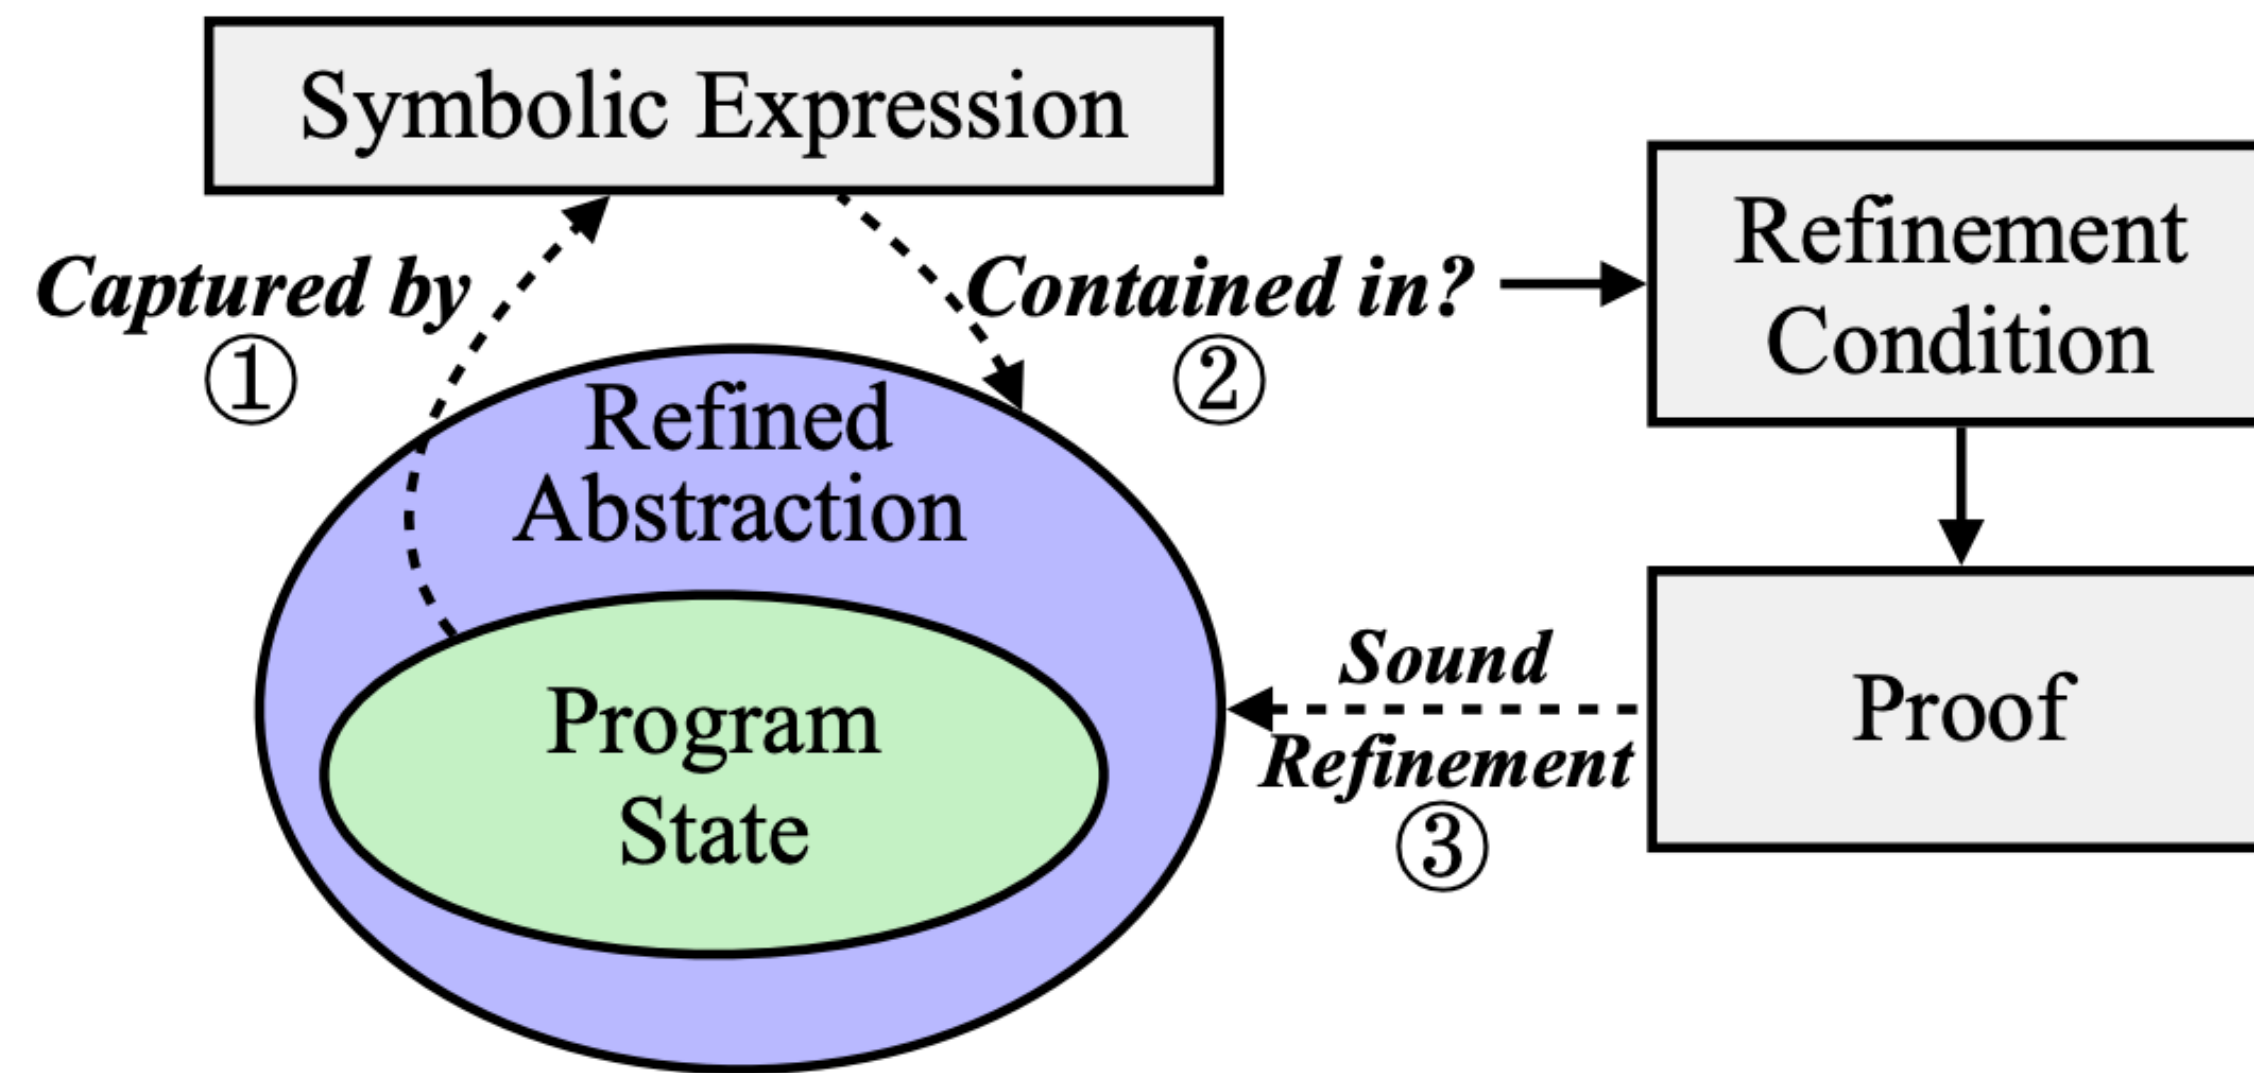
Keep the verifier simple, delegate nontrivial reasoning, and bridge the gap with proofs.

# Proof-Guided Abstraction Refinement

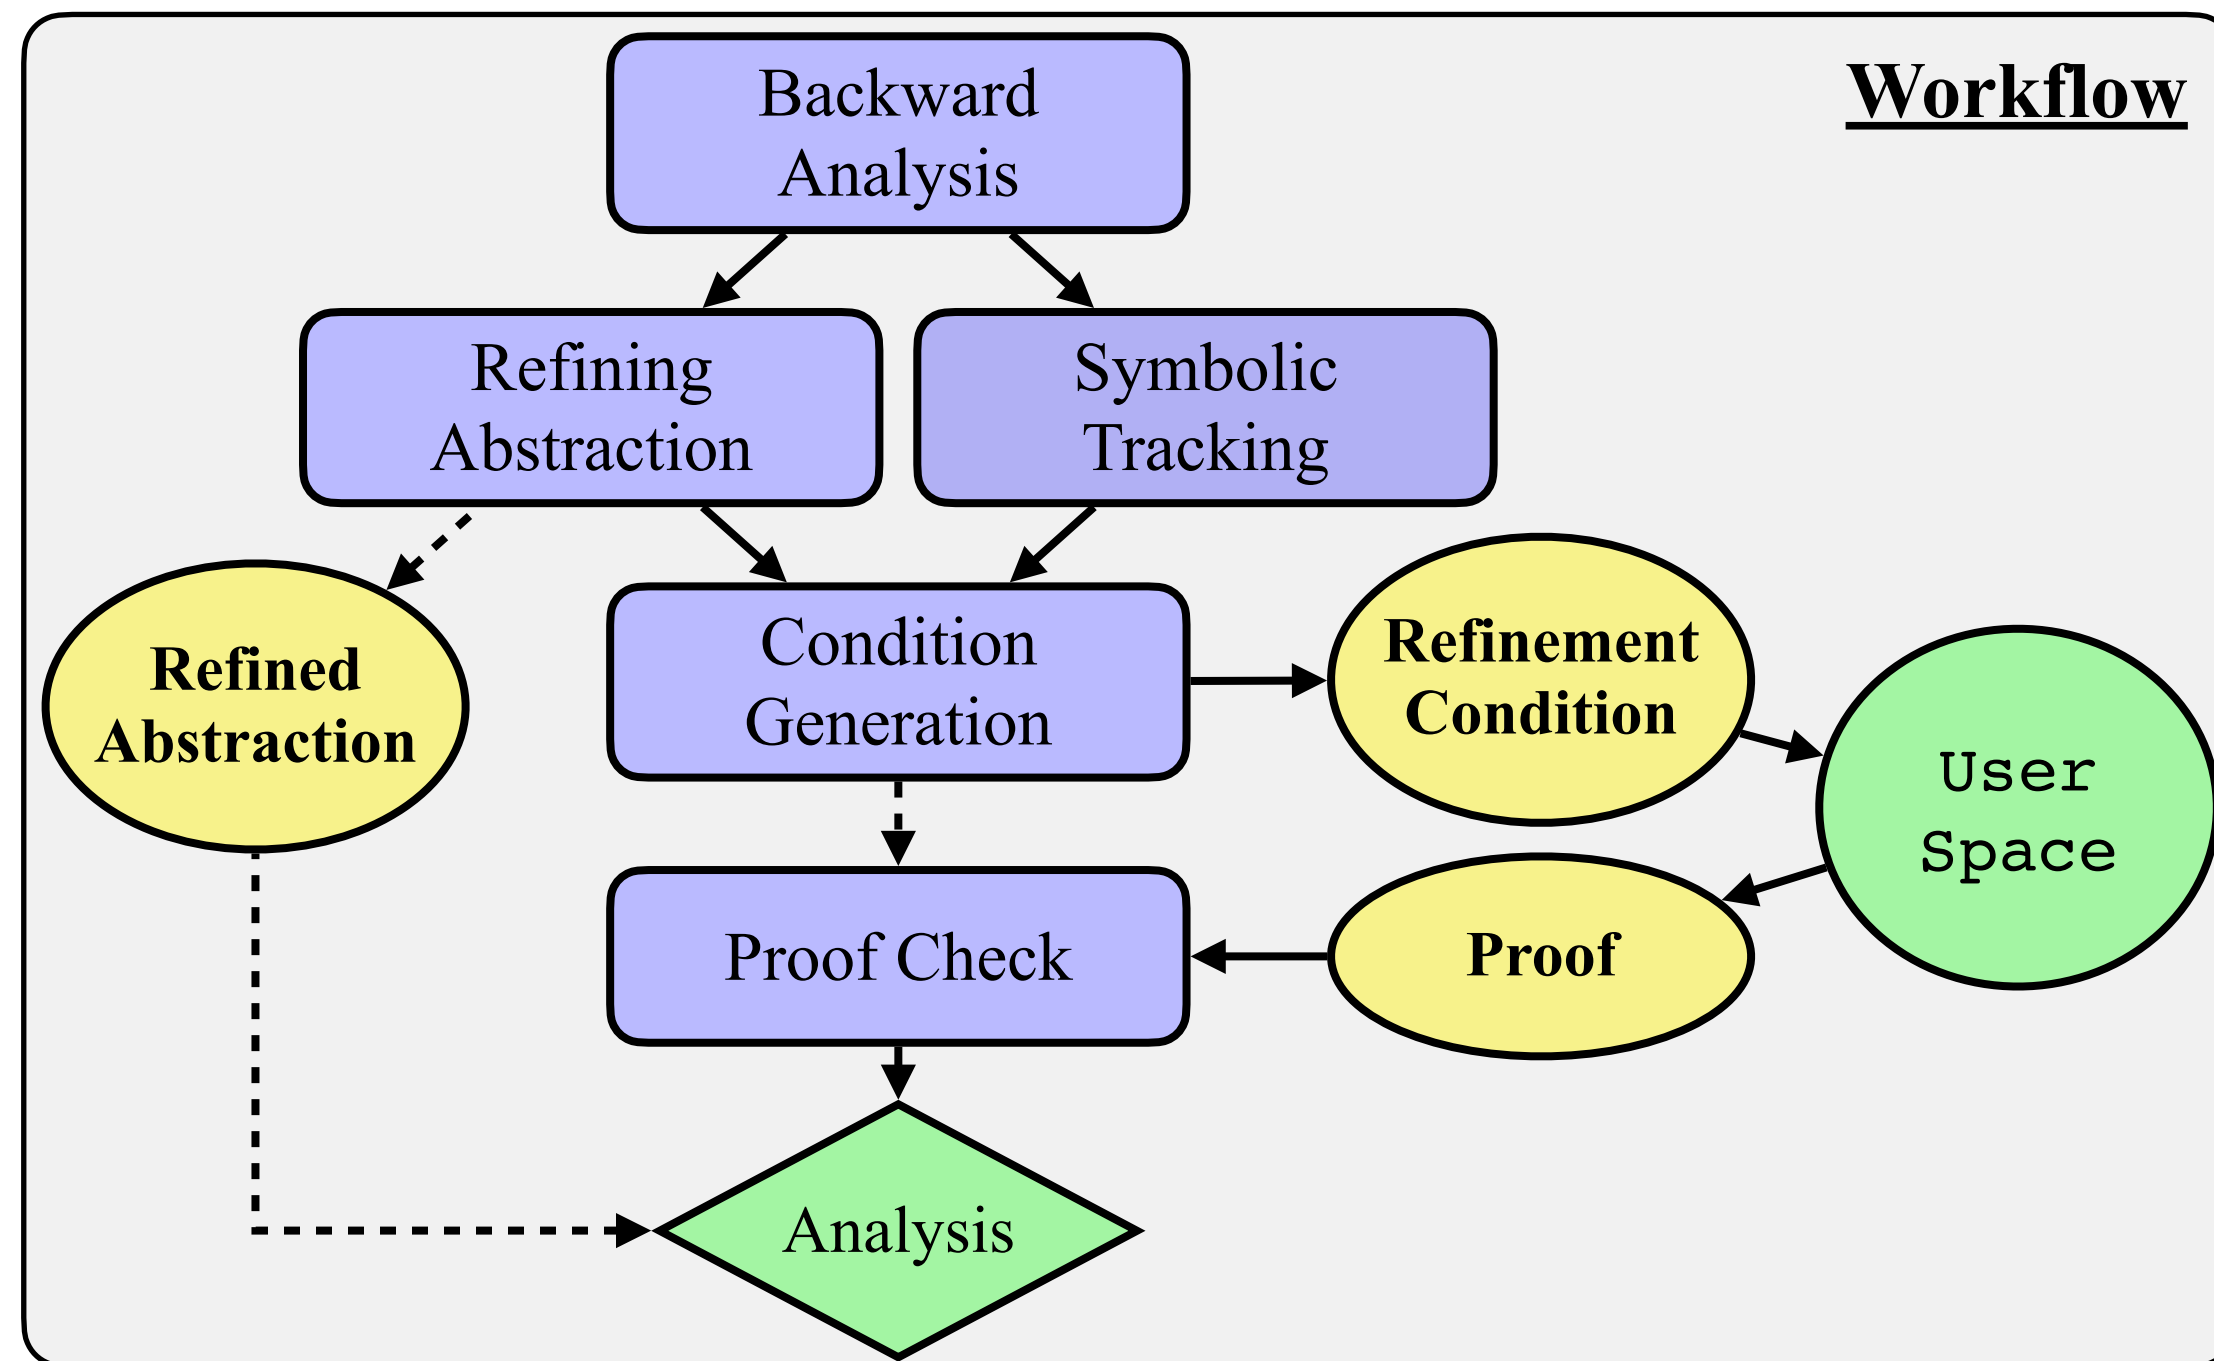


- On demand **abstraction refinement**:
  - The kernel verifier runs cheap analysis
  - Stops only when precision limits are reached
- Refinement soundness as **proof obligation**:
  - Encode refinement soundness as a formal formula
  - *User space* solver produces a formal proof
- Kernel space **proof checking**:
  - *Kernel proof checker* validates the proof
  - Verifier continues with the refined abstraction

# Proof-Guided Abstraction Refinement



# Workflow



```
... // instructions before are irrelevant
// (analysis ends, the suffix found)
r1 = map_lookup(...) // {} (r1 defined)
r2 = load_ctx(...) // {r1} (r2 defined)
r2 &= 0xf // {r1, r2}
r1 += r2 // {r1, r2}
r3 = 0xf - r2 // {r1, r2} (r3 defined, r2 added)
r1 += r3 // {r1, r3} (r3 added to the set)
r0 = *(u8*)r1 // {r1} (backward analysis starts)
```

**Listing 4.** Backward analysis to pinpoint the start location.

# Refinement

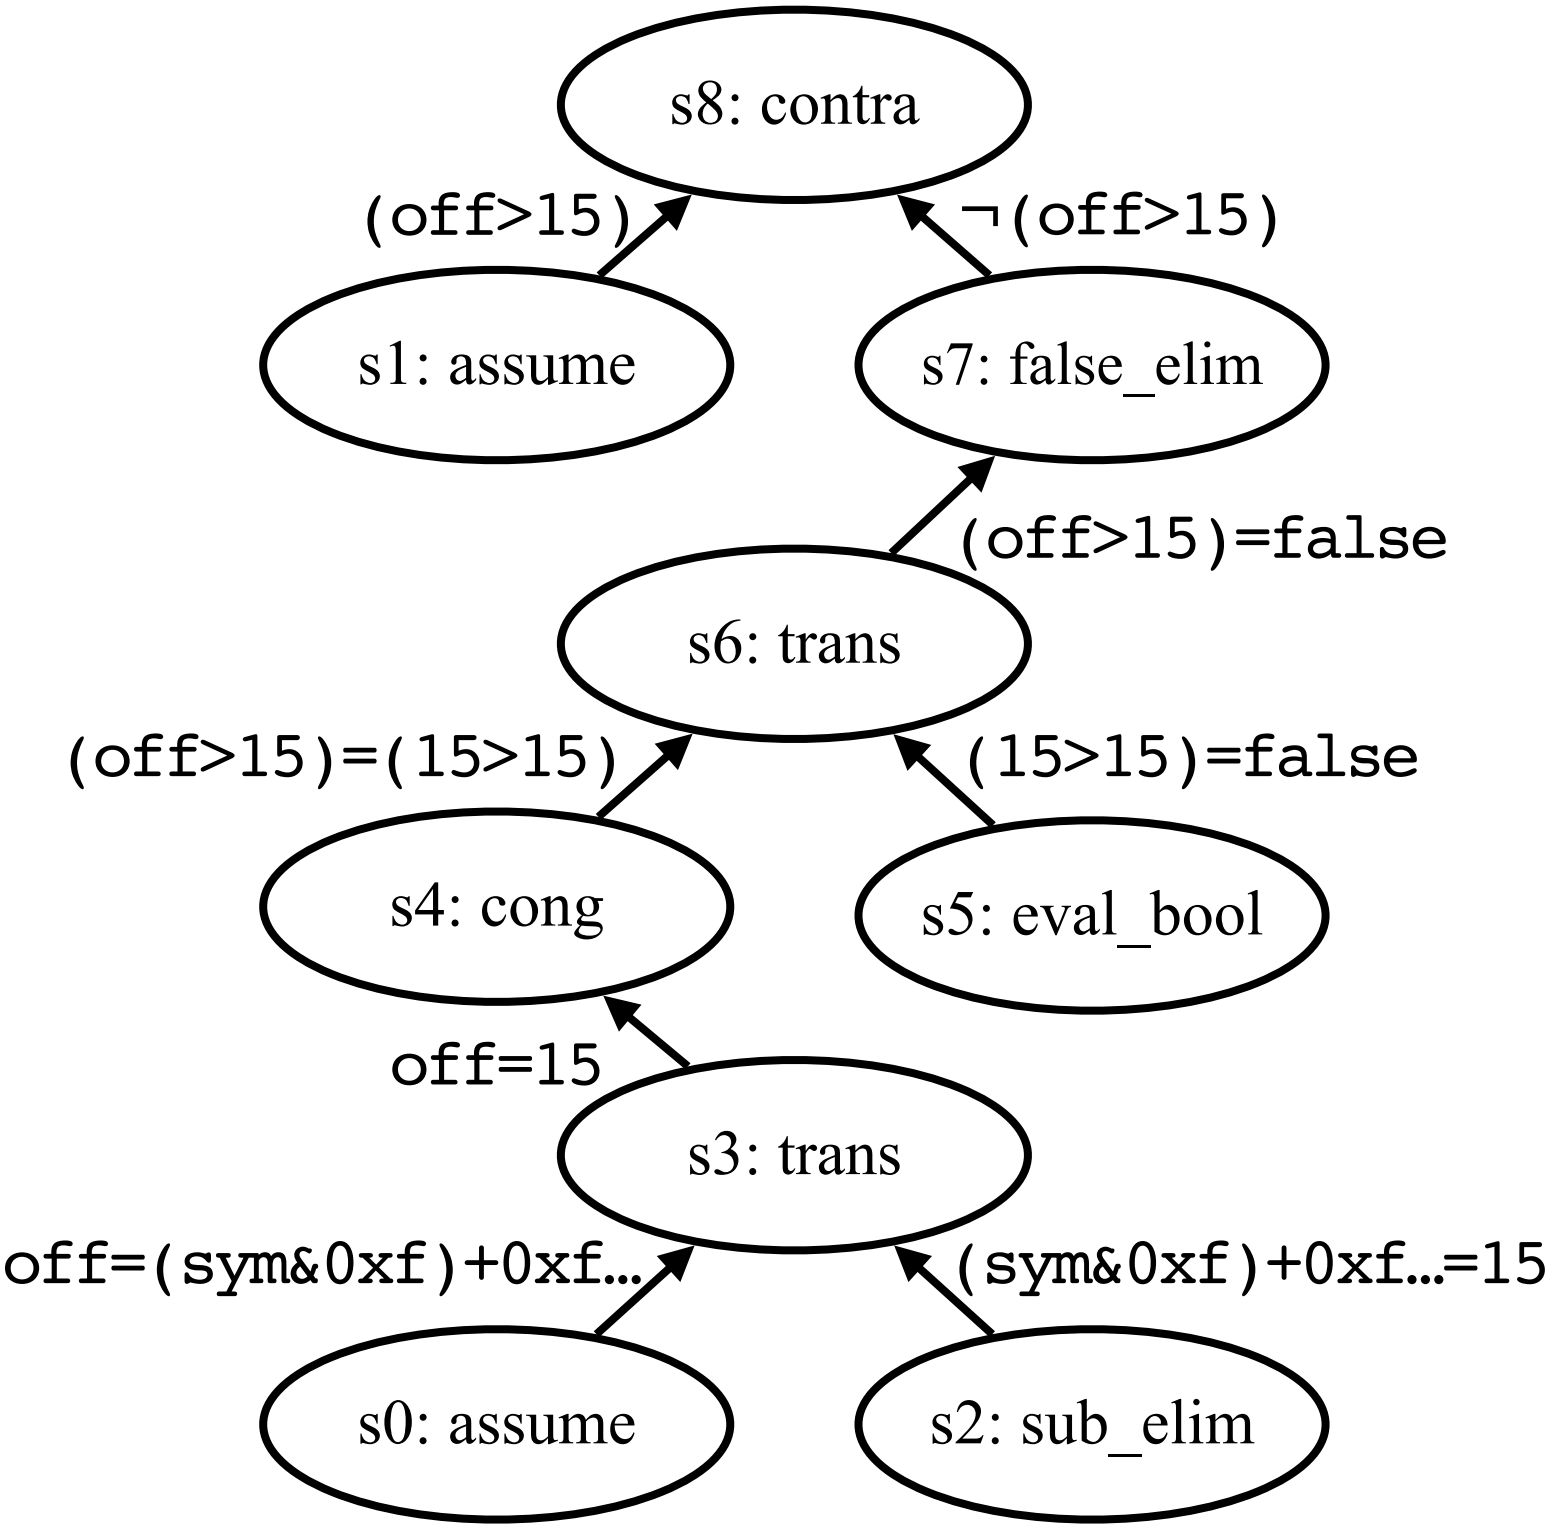
Instruction	Abstraction	Refinement
1: r1 = map_lookup(...)	$r1 = ptr(\underline{\text{size}=16})$	$r2 = sym$
2: r2 &= 0xf	$r2 \in [0, 15]$	$r2 = sym \& 0xf$
3: r1 += r2	$r1.off \in [0, 15]$	$r1.off = sym \& 0xf$
4: r3 = 0xf - r2	$r3 \in [0, 15]$	$r3 = 0xf - sym \& 0xf$
5: r1 += r3	$r1.off \in [0, 30]$	$r1.off = (sym \& 0xf) + (0xf - sym \& 0xf)$
6: r0 = *(u8*)r1	Unsafe, refine...	Condition: $r1.off \leq 15$
6: ...	$r1.off \in [0, 15]$	Proved and $r1.off$ <u>refined</u>



# Proof Check

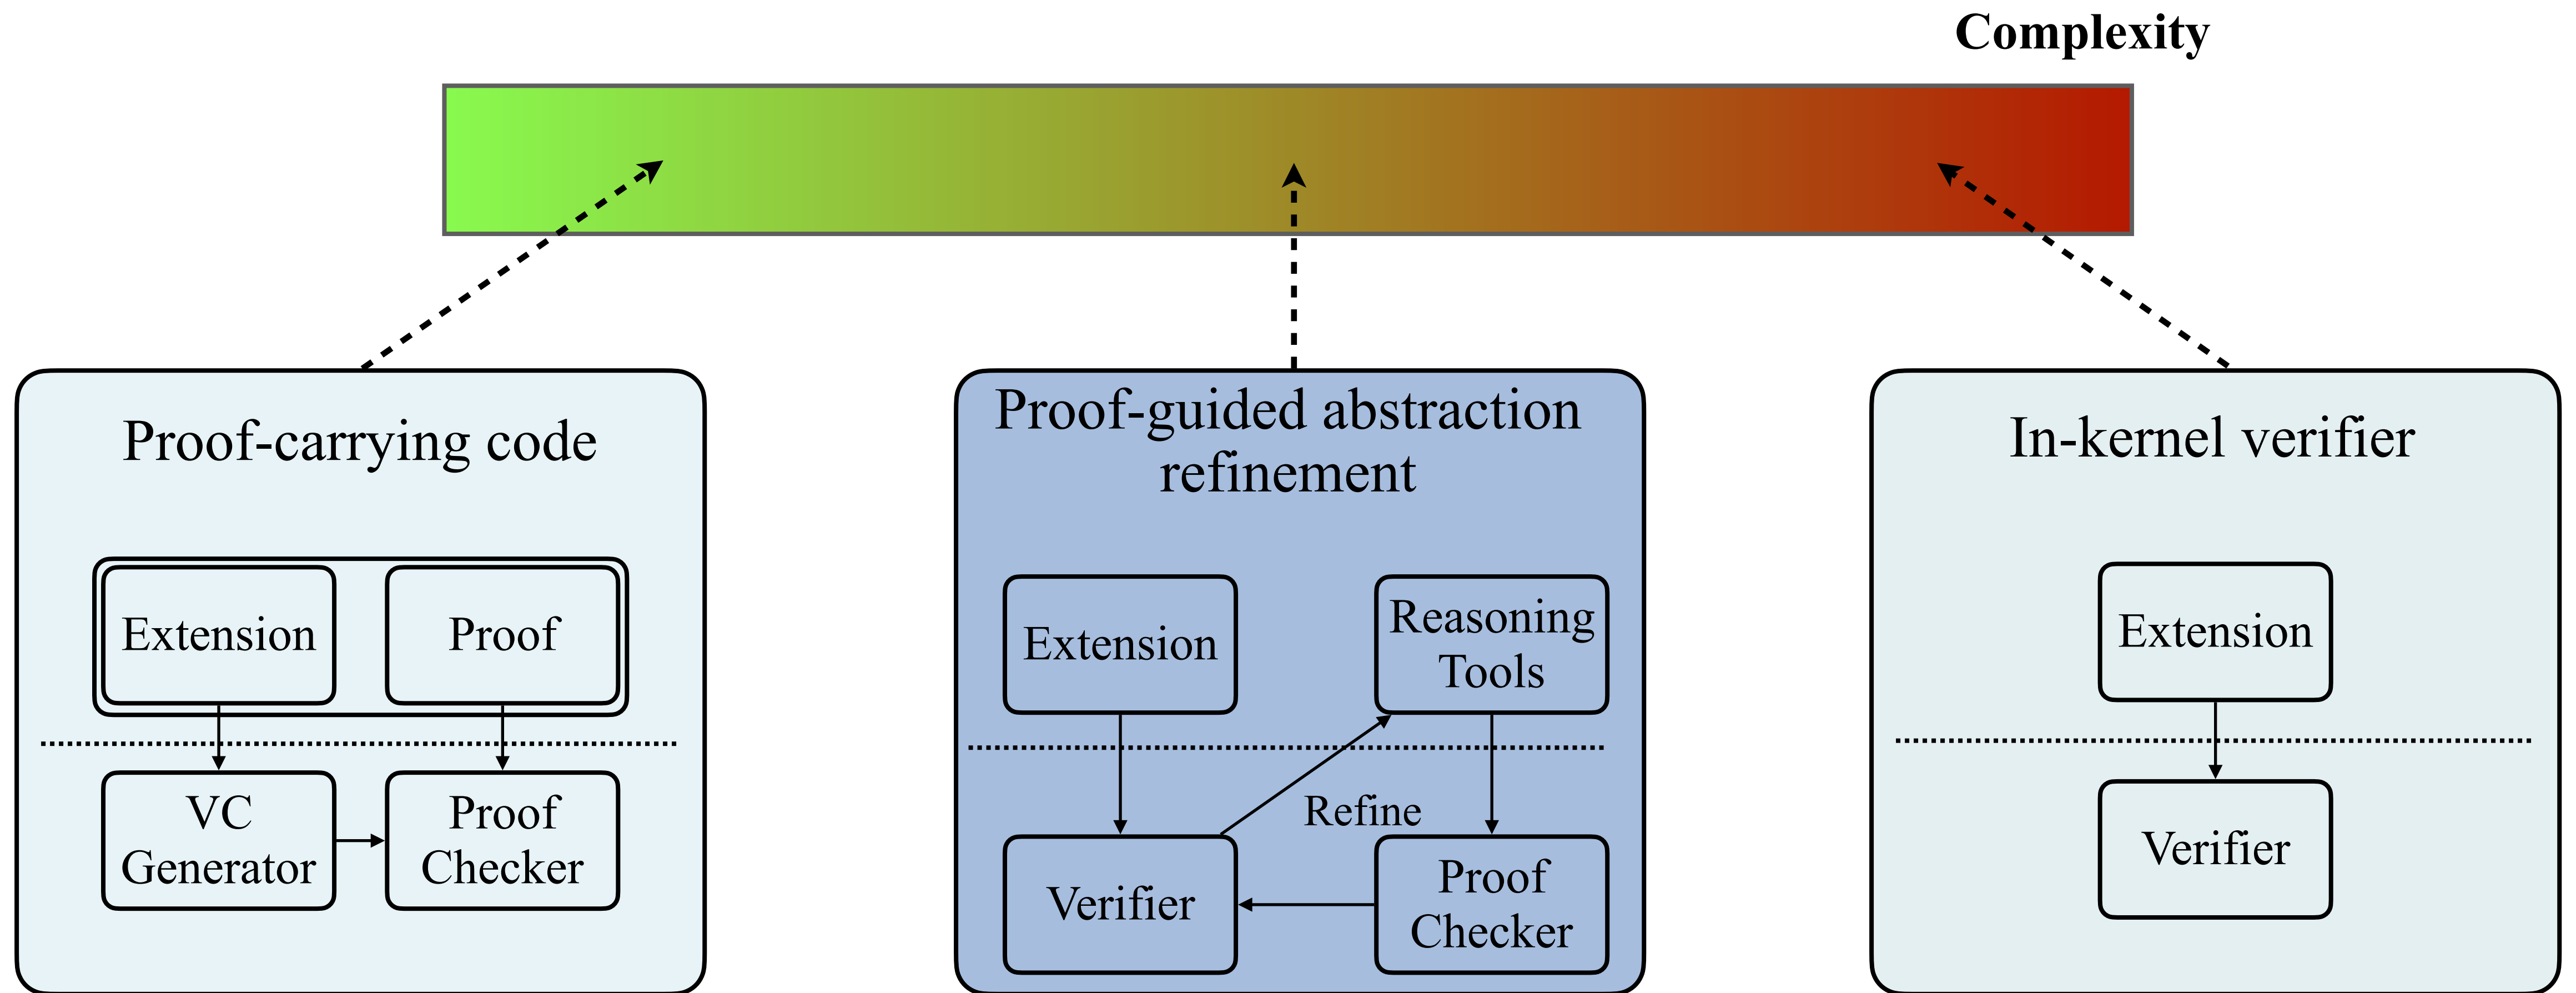
Step	Rule	Premise	Conclusion
<u>s0</u>	assume	-	off = (sym&0xf) + 0xf - (sym&0xf)
<u>s1</u>	assume	-	off > 15
<u>s2</u>	sub_elim	-	(sym&0xf) + 0xf - (sym&0xf) = 15
<u>s3</u>	trans	<u>s0,s2</u>	off = 15
<u>s4</u>	cong	<u>s3</u>	(off > 15) = (15 > 15)
<u>s5</u>	eval_bool	-	(15 > 15) = false
s6	trans	<u>s4,s5</u>	(off > 15) = false
<u>s7</u>	false_elim	s6	¬(off > 15)
s8	contra	<u>s1,s7</u>	<b>FALSE</b>

Rules:	$\frac{t_1 = t_2, \dots, t_{n-1} = t_n}{t_1 = t_n} \text{ trans}$	$\frac{t_1 = s_1, \dots, t_n = s_n}{f(t_1, \dots, t_n) = f(s_1, \dots, s_n)} \text{ cong}$
	$\frac{}{a + b - a = b} \text{ sub\_elim}$	$\frac{F = \text{false}}{\neg F} \text{ false\_elim} \quad \frac{F, \neg F}{\text{false}} \text{ contra}$



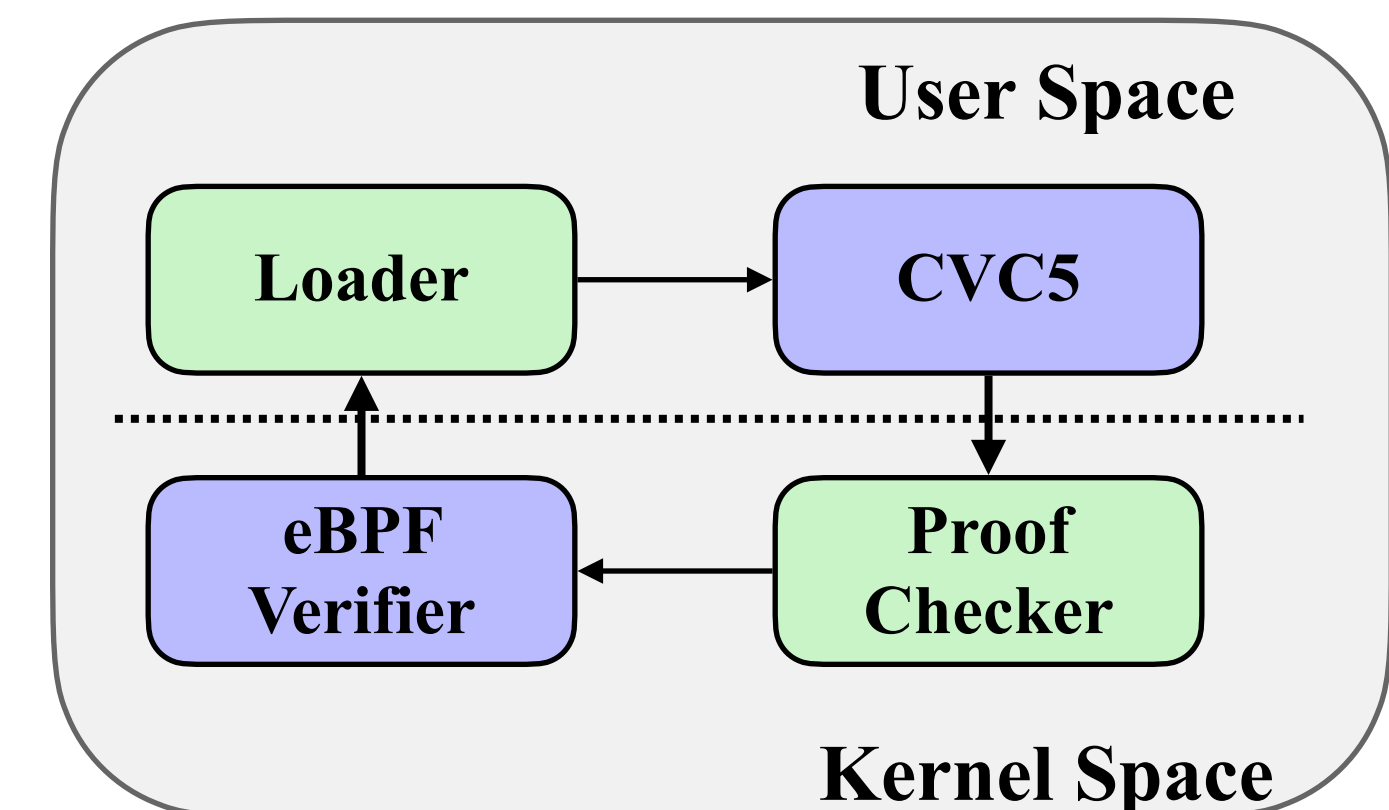



# Complexity





# BCF: eBPF Certificate Framework

- Refinement procedure in the verifier (1.7K LOC)
  - Triggered when analysis stalls
  - Produces a refinement condition
- Loader program
  - Receive the condition, translates it for the solver
- CVC5 SMT solver
  - Performs reasoning and generates a formal proof
- In-kernel proof checker (5K LOC)
  - Performs linear-time proof validation



# BCF: eBPF Certificate Framework

- Initial prototype open-sourced:
  - <https://github.com/SunHao-0/BCF/tree/artifact-evaluation>
- Continuous improvement:
  - <https://github.com/SunHao-0/BCF/tree/main>
  - Proof checker completely rewritten
  - Supports 50 proof rules:
    - 14 core, 33 boolean, 3 bitvector rules
  - 151 rewrite rules automatically converted from CVC5 RARE rewrites
- Ultimate goal: make this happen in the Linux kernel.



# Evaluation

- Compiler-driven approach to derive a dataset
- For the same program, compile it with different configurations
- For the same source, rejection in an alternative configuration indicates a false rejection
- 512 programs collected in total
  - from real-world project
  - compiled with widely-adopted compilers
- Sizes range from 0.5 to 376 KiB

Object	Project	Size	Loc	Description
bpf_lxc.o [30]	Cilium	269 KiB	2,450	Container identity and policy
bpf_host.o [29]	Cilium	376 KiB	2,086	Host level policy and route
pping_kern.o [91]	xdp-project	19 KiB	1,546	XDP packet timestamping
xdp_synproxy.o [92]	xdp-project	9.9 KiB	821	XDP-based SYN proxy
felix_bin_bpf.o [26]	Calico	188 KiB	2,162	Pod network policy
ksnoop.bpf.o [53]	BCC	6.3 KiB	461	Kernel function tracing




# Key Results

- 403 out of 512 programs automatically loaded
- 78.7% accepted, 21.3% still rejected
- Reasons for remaining rejections:
  - Refinement not triggered for four programs
  - 82 were due to condition not satisfied
  - 23 were due to reaching the one million instruction limit
- The first two cases can be solved with further engineering efforts
- The last case requires better loop handling

Metric	Min	Avg	Max
Refinement Frequency	1	446	16,048
Symbolic Track Length	7	102	373
Condition Size (bytes)	88	836	2,128
Proof Check Time ( $\mu$ s)	31	49	1,845
Proof Size (bytes)	136	541	46,296

# Key Results

- 403 out of 512 programs automatically loaded
- 78.7% accepted, 21.3% still rejected
- Reasons for remaining rejections:
  - Refinement not triggered for four programs
  - 82 were due to condition not satisfied
  - 23 were due to reaching the instruction limit
- The first two cases can be solved with further engineering
- The last case requires better loop handling





# Thank you!

